

A Comprehensive Review of *Game-Based RL*: Using Digital Games as Multi-Agent Reinforcement Learning Environments

Sergei Betchkov¹, Ziyi Wang^{1†}, Licheng Liu²

1. EigenLife ML Lab, Northeast Electric Power University, Jilin, China; 2. GeepSeek Team.

Correspondence to: dev.projectamadeus@outlook.com

†: Corresponding Author

Received: 2026/1/10, Revised: –, Accepted: 2026/2/20, Published: 2026/2/26

Abstract

The intersection of digital games and reinforcement learning (RL) has evolved from single-agent challenges into a dominant paradigm for studying multi-agent intelligence. This comprehensive review systematically examines the field of Game-Based Multi-Agent Reinforcement Learning (Game-MARL), positioning digital games as the premier testbeds for developing and evaluating algorithms that require coordination, competition, and strategic interaction among multiple learning agents. We first establish a taxonomy of multi-agent game environments, categorizing them by complexity (from cooperative to mixed-motivation scenarios), information structure (fully vs. partially observable), and communicative capacity. The review then traces the evolutionary trajectory of Game-MARL, from early zero-sum games to recent breakthroughs in complex team-play environments such as StarCraft II, Dota 2, and emerging open-world multiplayer sandboxes. We provide a comprehensive catalog of standardized benchmarks and platforms specifically designed for multi-agent research, analyzing their unique properties and suitability for different research questions. Core algorithmic challenges—including credit assignment, non-stationarity, scalability, and opponent modeling—are discussed within the context of game environments. Finally, we identify open problems and promising future directions, particularly the integration of large language models for human-AI collaboration and the potential for procedurally generated game worlds to test generalization in multi-agent systems. This review aims to serve as a foundational resource for researchers entering the field and a structured reference for those advancing the frontier of multi-agent intelligence through games.

Keywords: Game-Based Reinforcement Learning, Digital Game Environments, Deep Reinforcement Learning, Multi-Agents

Contents

1	Introduction	4
2	Preliminaries	4
2.1	Single-Agent Reinforcement Learning	4
2.2	Multi-Agent Reinforcement Learning Formalism	5
2.3	Foundational Challenges in Multi-Agent Reinforcement Learning	6
3	A Taxonomy of Multi-Agent Game Environments	7
3.1	Classification by Game Genre	7
3.1.1	Board Games and Classical Games	7
3.1.2	Real-Time Strategy Games	7
3.1.3	Multiplayer Online Battle Arena Games	8
3.1.4	First-Person Shooter Games	8
3.1.5	Card Games	8
3.1.6	Open-World Sandbox Games	8
3.2	Classification by Interaction Type	9
3.2.1	Fully Cooperative Environments	9
3.2.2	Fully Competitive Environments	9
3.2.3	Mixed-Motive Environments	9
3.3	Classification by Information Structure	9
3.3.1	Perfect Information Environments	10
3.3.2	Imperfect Information Environments	10
3.4	Classification by Scale	10

4	Core Algorithms and Methodologies	10
4.1	Value Decomposition Methods	10
4.1.1	Foundations of Value Decomposition	10
4.1.2	Value Decomposition Networks	11
4.1.3	QMIX and Monotonic Factorization	11
4.1.4	Advanced Value Decomposition Methods	11
4.1.5	Theoretical Insights and Limitations	12
4.2	Centralized Training with Decentralized Execution	12
4.2.1	Foundations of CTDE	12
4.2.2	MADDPG	12
4.2.3	COMA	13
4.2.4	MAPPO	13
4.2.5	Comparison of CTDE Methods	13
4.3	Communication-Based Methods	13
4.3.1	Learning to Communicate	14
4.3.2	Structured Communication Protocols	14
4.3.3	Emergent Communication and Language	14
4.3.4	Challenges in Communication-Based Methods	15
4.4	League Training and Population-Based Methods	15
4.4.1	Self-Play and Its Limitations	15
4.4.2	AlphaStar League Training	15
4.4.3	OpenAI Five and Prioritized Training	15
4.4.4	PSRO and Game-Theoretic Methods	15
4.4.5	Population-Based Training	16
4.5	Algorithms for Mixed-Motive Environments	16
4.5.1	Social Dilemmas and Reciprocity	16
4.5.2	Opponent Modeling and Intention Recognition	16
4.5.3	Social Value Alignment	16
4.6	Algorithmic Frameworks and Design Patterns	16
4.6.1	Actor-Critic Architectures	16
4.6.2	Attention and Transformer Mechanisms	17
4.6.3	Hierarchical and Temporal Abstraction	17
5	Key Applications and Milestones	17
5.1	Classic Board Games: Foundations of Multi-Agent Reasoning	17
5.1.1	Chess: Deep Blue and Beyond	17
5.1.2	Go: AlphaGo and AlphaZero	17
5.1.3	Multi-Player Board Games: Diplomacy and Hanabi	18
5.2	Real-Time Strategy: StarCraft and the Challenge of Macro-Strategy	18
5.2.1	The StarCraft AI Competition	18
5.2.2	AlphaStar: Grandmaster Level in StarCraft II	18
5.2.3	SMAC: A Focused Benchmark for Cooperative Learning	19
5.3	Multiplayer Online Battle Arena: Dota and Team Coordination	19
5.3.1	OpenAI Five: Defeating World Champions in Dota 2	19
5.3.2	JueWu: Industrial-Scale Multi-Agent Learning	19
5.3.3	Coordination Mechanisms in MOBA Games	20
5.4	First-Person Shooters: Real-Time Coordination and Spatial Reasoning	20
5.4.1	ViZDoom and Deep Learning from Pixels	20
5.4.2	DeepMind Lab and 3D Navigation	20
5.4.3	Counter-Strike: Modern Benchmarks	20
5.5	Open-World Sandboxes: Minecraft and Open-Ended Learning	20
5.5.1	Project Malmö and the MineRL Competition	20
5.5.2	Voyager: Open-Ended Exploration with Large Language Models	21
5.5.3	Multi-Agent Minecraft: Coordination and Collaboration	21
5.6	Card Games: Imperfect Information and Probabilistic Reasoning	21
5.6.1	Poker: Pluribus and Equilibrium Finding	21
5.6.2	Hearthstone and Magic: The Gathering	21
5.7	Lessons Learned and Cross-Cutting Insights	21

6	Challenges and Open Problems	22
6.1	Sample Efficiency	22
6.1.1	Magnitude of the Problem	22
6.1.2	Sources of Inefficiency	22
6.1.3	Promising Directions	22
6.1.4	Open Questions	23
6.2	Generalization	23
6.2.1	Types of Generalization	23
6.2.2	The Generalization Gap	23
6.2.3	Sources of Poor Generalization	23
6.2.4	Approaches to Improving Generalization	24
6.2.5	Open Questions	24
6.3	Scalability	24
6.3.1	Curse of Dimensionality	24
6.3.2	Computational Scaling	24
6.3.3	Many-Agent Challenges	24
6.3.4	Promising Directions	25
6.3.5	Open Questions	25
6.4	Credit Assignment	25
6.4.1	Temporal Credit Assignment	25
6.4.2	Spatial Credit Assignment	25
6.4.3	Structural Credit Assignment	26
6.4.4	Open Questions	26
6.5	Non-Stationarity and Convergence	26
6.5.1	The Non-Stationarity Problem	26
6.5.2	Convergence Guarantees	26
6.5.3	Practical Approaches	26
6.5.4	Open Questions	27
6.6	Opponent Modeling and Theory of Mind	27
6.6.1	The Recursion Problem	27
6.6.2	Computational Challenges	27
6.6.3	Open Questions	27
6.7	Explainability and Interpretability	27
6.7.1	The Black Box Problem	27
6.7.2	Explainability Methods	28
6.7.3	Open Questions	28
6.8	Safety and Robustness	28
6.8.1	Distributional Shift and Robustness	28
6.8.2	Safe Exploration	28
6.8.3	Alignment and Value Learning	28
6.8.4	Open Questions	29
6.9	Evaluation and Benchmarking	29
6.9.1	Statistical Challenges	29
6.9.2	Benchmark Design	29
6.9.3	Reproducibility	29
6.9.4	Open Questions	29
7	Conclusions and Future Directions	29
7.1	Summary of Key Insights	30
7.1.1	The Co-Evolution of Environments and Algorithms	30
7.1.2	The Primacy of Coordination Mechanisms	30
7.1.3	The Power of Population-Based Training	30
7.1.4	The Centrality of Representation	30
7.2	Future Directions	30
7.2.1	Foundation Models for Multi-Agent Systems	31
7.2.2	Open-Ended Learning and Auto-Curricula	31
7.2.3	Human-AI Coordination and Teaming	31
7.2.4	Theoretical Foundations	31
7.2.5	Beyond Games: Real-World Transfer	31
7.2.6	Social and Ethical Implications	32
7.3	Concluding Remarks	32

1 Introduction

The quest for artificial general intelligence (AGI) has long been intertwined with the development of agents capable of operating in complex, dynamic, and uncertain environments. Among the various testbeds proposed over the decades, digital games have emerged as one of the most fertile and influential platforms for artificial intelligence research. Unlike real-world physical environments, which are often costly, slow, and difficult to replicate, digital games offer a unique combination of high-fidelity simulation, precise controllability, and massive parallelizability (Bellemare et al., 2013). These properties have positioned games at the heart of modern reinforcement learning (RL), enabling a series of breakthroughs that have progressively redefined the boundaries of machine intelligence.

The historical trajectory of game-based AI research reveals a consistent pattern: each leap in algorithmic capability has been accompanied by a corresponding increase in the complexity of the game environments used for validation. The early success of deep Q-networks (DQN) on the Atari 2600 benchmark demonstrated that deep neural networks could learn directly from high-dimensional sensory inputs to achieve human-level performance in discrete action spaces (Mnih et al., 2013, 2015). This milestone catalyzed a wave of research that rapidly expanded into more challenging domains. Subsequently, the mastery of Go through AlphaZero showcased the power of combining deep reinforcement learning with Monte Carlo tree search in perfect-information, two-player zero-sum games (Silver et al., 2017). These achievements, while remarkable, primarily addressed single-agent scenarios or two-player adversarial settings with simple interaction structures.

The contemporary frontier of game-based RL, however, is increasingly defined by the transition from single-agent to multi-agent systems. Real-world intelligence rarely operates in isolation; from autonomous driving and financial markets to robotics swarms and social reasoning, agents must continuously interact, negotiate, compete, and collaborate with other intelligent entities. This shift in focus has positioned multi-agent reinforcement learning (MARL) as a central pillar of modern AI research (Hernandez-Leal et al., 2019). Digital games provide an ideal microscope for examining these complex interactions. Environments such as *StarCraft II* offer partially observable, large-scale military confrontations requiring hierarchical coordination (Vinyals et al., 2019; Samvelyan et al., 2019), while titles like *Dota 2* demand real-time team play with long-term strategic horizons (Bernier et al., 2019). More recently, open-world sandboxes like *Minecraft* have emerged as platforms for studying open-ended collaboration and human-AI interaction (Guss et al., 2019; Wang et al., 2023).

Despite the rapid proliferation of research at the intersection of games and MARL, the field lacks a structured, comprehensive synthesis that connects the evolution of game environments to the development of multi-agent algorithms. Existing surveys often focus either on RL algorithms in isolation or on specific game genres without adequately addressing the unique challenges posed by multi-agent dynamics. This article aims to fill that gap by providing a comprehensive review of Game-Based Multi-Agent Reinforcement Learning (Game-MARL). We systematically examine the landscape of digital game environments designed for MARL research, propose a taxonomy for classifying these environments, trace the key algorithmic innovations they have inspired, and discuss the persistent challenges that define the current research frontier. By doing so, we seek to provide researchers—both newcomers and veterans—with a structured roadmap of a field that stands at the crossroads of game design, multi-agent systems, and artificial general intelligence.

2 Preliminaries

Before delving into the complex interplay between digital games and multi-agent reinforcement learning, it is essential to establish the foundational concepts upon which the field is built. This section provides a systematic overview of the formal frameworks and core challenges that underpin multi-agent reinforcement learning in game environments. We begin with a concise review of single-agent reinforcement learning, then extend this formalism to the multi-agent setting through the framework of Markov games, and finally discuss the fundamental challenges that distinguish multi-agent learning from its single-agent counterpart.

2.1 Single-Agent Reinforcement Learning

Reinforcement learning is a paradigm in which an agent learns to make sequential decisions by interacting with an environment. The standard mathematical formalism for single-agent reinforcement learning is the Markov decision process, which provides a framework for modeling decision-making in stochastic environments where outcomes are partly random and partly under the control of the decision-maker (Sutton and Barto, 1998). A Markov decision process is defined by the tuple $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$, where $\mathcal{S}$ represents the set of states, $\mathcal{A}$ denotes the set of actions available to the agent, $\mathcal{P} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$ is the transition probability function defining the dynamics of the environment, $\mathcal{R} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ is the reward function that provides feedback to the agent, and $\gamma \in [0, 1]$ is a discount factor that balances immediate and future rewards.

At each discrete time step t , the agent observes the current state $S_t \in \mathcal{S}$, selects an action $A_t \in \mathcal{A}$ according to its policy $\pi : \mathcal{S} \rightarrow \mathcal{P}(\mathcal{A})$, which maps states to probability distributions over actions. The environment then transitions to a new state $S_{t+1} \sim \mathcal{P}(\cdot | S_t, A_t)$ and provides a scalar reward $R_{t+1} = \mathcal{R}(S_t, A_t, S_{t+1})$. The fundamental objective of the agent is to learn a policy that maximizes the expected discounted cumulative return, formally expressed as

$$G_t = \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \right]. \quad (1)$$

A central concept in reinforcement learning is the value function, which estimates the expected return from a given state or state-action pair. The state-value function under policy π is defined as $V_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s]$, while the action-value function $Q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$ provides the expected return for taking a specific action in a given state and subsequently following policy π . These value functions satisfy fundamental recursive relationships known as the Bellman equations (Bellman, 1957), which form the theoretical foundation for numerous reinforcement learning algorithms. For the action-value function, the Bellman equation is given by

$$Q_\pi(s, a) = \sum_{s' \in \mathcal{S}} \mathcal{P}(s' | s, a) \left[\mathcal{R}(s, a, s') + \gamma \sum_{a' \in \mathcal{A}} \pi(a' | s') Q_\pi(s', a') \right]. \quad (2)$$

The optimal action-value function $Q^*(s, a)$ satisfies the Bellman optimality equation, which eliminates the dependence on a specific policy by taking the maximum over actions:

$$Q^*(s, a) = \sum_{s' \in \mathcal{S}} \mathcal{P}(s' | s, a) \left[\mathcal{R}(s, a, s') + \gamma \max_{a' \in \mathcal{A}} Q^*(s', a') \right]. \quad (3)$$

This equation forms the basis for value-based reinforcement learning algorithms such as Q-learning (Watkins and Dayan, 1992), which iteratively updates Q-value estimates using the Bellman optimality equation as a target.

An alternative approach to reinforcement learning is policy-based methods, which directly parameterize the policy $\pi_\theta(a|s)$ with parameters θ and optimize the expected return using gradient ascent. The policy gradient theorem (Sutton et al., 2000) provides the theoretical foundation for these methods, deriving an expression for the gradient of the performance objective $J(\theta)$ with respect to the policy parameters:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(a|s) Q^{\pi_\theta}(s, a)]. \quad (4)$$

Modern deep reinforcement learning integrates these classical concepts with deep neural networks as function approximators, enabling agents to operate in high-dimensional state spaces such as raw pixel inputs from game screens. The deep Q-network algorithm represents a seminal advancement in this direction, demonstrating that reinforcement learning could achieve human-level performance on a range of Atari 2600 games by combining Q-learning with convolutional neural networks (Mnih et al., 2013, 2015). This breakthrough catalyzed extensive research into deep reinforcement learning methods, including policy gradient approaches such as the proximal policy optimization algorithm (Schulman et al., 2017), which has become a standard choice for many game-based reinforcement learning applications due to its stability and sample efficiency.

2.2 Multi-Agent Reinforcement Learning Formalism

When multiple autonomous agents coexist and interact within a shared environment, the single-agent Markov decision process framework proves insufficient, as the dynamics become non-stationary from the perspective of any individual agent due to the evolving policies of other agents. The natural extension of Markov decision processes to multi-agent settings is the stochastic game, also known as the Markov game, originally introduced by (Shapley, 1953) and later formalized for reinforcement learning by (Littman, 1994). A stochastic game provides a mathematical framework for modeling multi-agent sequential decision-making in which multiple agents interact over time, each pursuing their own objectives.

A stochastic game is defined by the tuple $(\mathcal{N}, \mathcal{S}, \{\mathcal{A}^i\}_{i \in \mathcal{N}}, \mathcal{P}, \{\mathcal{R}^i\}_{i \in \mathcal{N}}, \gamma)$, where $\mathcal{N} = \{1, 2, \dots, n\}$ represents the set of agents, $\mathcal{S}$ is the state space shared by all agents, and $\mathcal{A}^i$ denotes the action space available to agent i . The joint action space is the Cartesian product $\mathcal{A} = \times_{i \in \mathcal{N}} \mathcal{A}^i$, and at each time step, agents simultaneously select actions forming a joint action $\mathbf{a} = (a^1, a^2, \dots, a^n) \in \mathcal{A}$. The transition probability function $\mathcal{P} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$ governs the environment dynamics based on the current state and the joint action of all agents. Each agent i receives its own reward according to its reward function $\mathcal{R}^i : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$, and all agents share a common discount factor $\gamma \in [0, 1)$.

The behavior of each agent is governed by its individual policy $\pi^i : \mathcal{S} \rightarrow \mathcal{P}(\mathcal{A}^i)$, which maps states to probability distributions over its own action space. The collection of all agent policies forms a joint policy $\pi = (\pi^1, \pi^2, \dots, \pi^n)$. Each agent aims to maximize its own expected discounted cumulative return, which now depends on the joint policy of all agents:

$$G_t^i = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k \mathcal{R}_{t+k+1}^i \right]. \quad (5)$$

This formulation captures a wide spectrum of multi-agent interaction scenarios that are particularly relevant to digital games. When all agents share identical reward functions such that $\mathcal{R}^i = \mathcal{R}^j$ for all $i, j \in \mathcal{N}$, the setting reduces to a fully cooperative stochastic game, exemplified by team-based game modes where agents must coordinate to achieve common objectives. At the opposite extreme, zero-sum stochastic games, where $\sum_{i \in \mathcal{N}} \mathcal{R}^i = 0$, model purely competitive interactions such as those found in two-player adversarial games like chess or Go. Between these poles lie general-sum stochastic games, which capture mixed-motive interactions characteristic of many modern multiplayer games where agents may have partially aligned and partially conflicting objectives (Busoniu et al., 2008).

The value functions in multi-agent settings must account for the policies of other agents. For a given agent i , the state-value function under joint policy π is defined as $V_\pi^i(s) = \mathbb{E}_\pi[G_t^i | S_t = s]$, and the action-value function $Q_\pi^i(s, \mathbf{a})$ gives the expected return for agent i when the joint action $\mathbf{a}$ is taken in state s and all agents subsequently follow π . For cooperative

settings where agents share a common reward function, it is often useful to define a joint action-value function $Q_\pi(s, \mathbf{a})$ that represents the expected return for the team. A key insight in modern multi-agent reinforcement learning is that this joint action-value function can be factorized into components associated with individual agents or subgroups. The value decomposition network approach (Sunehag et al., 2017) expresses this factorization as

$$Q_\pi(s, \mathbf{a}) = \sum_{i=1}^n Q_{\pi^i}^i(s, a^i), \quad (6)$$

where each $Q_{\pi^i}^i(s, a^i)$ depends only on the individual agent's action. This additive factorization enables scalable learning by avoiding the exponential blow-up of the joint action space. A more general factorization is provided by the QMIX algorithm (Rashid et al., 2018), which enforces monotonicity constraints through a mixing network:

$$Q_\pi(s, \mathbf{a}) = f_{\text{mix}}\left(s, Q_{\pi^1}^1(s, a^1), \dots, Q_{\pi^n}^n(s, a^n)\right), \quad (7)$$

with the constraint that $\frac{\partial Q_\pi}{\partial Q_{\pi^i}^i} \geq 0$ for all i , ensuring that the joint action-value function is monotonic in each individual utility.

2.3 Foundational Challenges in Multi-Agent Reinforcement Learning

The transition from single-agent to multi-agent reinforcement learning introduces a set of foundational challenges that fundamentally alter the nature of the learning problem. These challenges arise from the interactive and adaptive nature of multiple learning agents and have profound implications for algorithm design in game environments.

The first and perhaps most pervasive challenge is the non-stationarity of the environment from each agent's perspective. In single-agent reinforcement learning, the environment dynamics are stationary, meaning that the transition probabilities and reward functions remain constant throughout learning. This stationarity underpins the theoretical convergence guarantees of algorithms such as Q-learning. However, in multi-agent settings, the environment dynamics from the viewpoint of any individual agent depend on the policies of all other agents, which are themselves evolving concurrently through learning (Hernandez-Leal et al., 2019). As other agents improve and modify their behavior, the effective transition dynamics perceived by a focal agent change continuously, violating the Markov property that the environment is conditionally stationary given the agent's own actions. This non-stationarity poses significant challenges for learning convergence and necessitates algorithmic approaches that can cope with continually shifting dynamics.

A second fundamental challenge is the multi-agent credit assignment problem, which becomes particularly salient in cooperative game settings. When multiple agents collaborate to achieve a shared objective, the global reward signal received after a sequence of joint actions provides no inherent information about the contribution of each individual agent to the overall outcome (Chang et al., 2003). This credit assignment problem is exacerbated in game environments with delayed rewards, where the consequences of individual actions may only become apparent after many time steps. Determining which agents' actions were responsible for successful outcomes, and which may have contributed to failures, is essential for effective cooperative learning. Various approaches have been developed to address this challenge, including difference rewards (Wolpert and Tumer, 2001) and value decomposition methods represented by Equations (6) and (7).

Scalability represents a third major challenge, encompassing both the exponential growth of the joint action space and the computational demands of learning with many agents. In a game with n agents each having $|\mathcal{A}|$ actions, the joint action space grows as $|\mathcal{A}|^n$, rendering exhaustive exploration and representation infeasible for all but the smallest problems. This combinatorial explosion necessitates methods that can factorize across agents and exploit structural regularities in the interaction. Furthermore, as the number of agents increases, the variance of gradient estimates in policy gradient methods often grows, and the coordination requirements become increasingly complex (Papoudakis et al., 2020). The policy gradient theorem for multi-agent settings extends Equation (4) to account for the policies of other agents:

$$\nabla_{\theta^i} J(\theta^i) = \mathbb{E}_\pi \left[\nabla_{\theta^i} \log \pi_{\theta^i}^i(a^i | s) Q^\pi(s, \mathbf{a}) \right], \quad (8)$$

where the expectation is now taken over the joint policy of all agents, and the Q-function depends on the joint action.

Partial observability is another pervasive challenge in multi-agent game environments. While the stochastic game formalism assumes that all agents have access to the full state, many realistic game scenarios involve agents that only observe a local view of the environment. In first-person shooter games, for instance, each agent perceives only what is within their field of view, and in real-time strategy games, the classic fog of war conceals opponent movements and base expansions. This setting is formalized as a partially observable stochastic game, where each agent receives individual observations $o^i \in \mathcal{O}^i$ according to an observation function $\mathcal{O}^i : \mathcal{S} \rightarrow \mathcal{P}(\mathcal{O}^i)$ (Hansen et al., 2004). Partial observability introduces additional complexity, as agents must maintain internal beliefs or memory to infer the true state and the likely intentions of other agents. The optimal policy in such settings becomes a function of the observation history rather than the current state:

$$\pi^i : (\mathcal{O}^i)^* \rightarrow \mathcal{P}(\mathcal{A}^i), \quad (9)$$

mapping entire observation histories to action distributions.

Finally, opponent modeling emerges as both a necessity and a challenge in multi-agent learning. In competitive and mixed-motive game settings, an agent's optimal strategy depends critically on the strategies employed by opponents. The

ability to model and predict opponent behavior can provide significant advantages, enabling more effective counter-strategies and adaptation (Albrecht and Stone, 2018). A simple opponent modeling approach might maintain empirical estimates of opponent policies:

$$\hat{\pi}^{-i}(a^{-i}|s) \approx \frac{\sum_t \mathbb{I}[A_t^{-i} = a^{-i}, S_t = s]}{\sum_t \mathbb{I}[S_t = s]}, \quad (10)$$

where $\hat{\pi}^{-i}$ represents the estimated joint policy of all agents other than i . However, opponent modeling introduces a recursive reasoning problem, as each agent's model of others may itself be based on those others' models of itself. This recursive structure can lead to infinite regress, and practical approaches must impose tractable approximations. Moreover, in environments with multiple interacting agents, the computational burden of maintaining models for all other agents can become substantial.

These foundational challenges collectively define the research frontier of multi-agent reinforcement learning in game environments. Each challenge has spawned extensive research directions, and the algorithms discussed in subsequent sections of this review can be understood as attempts to address one or more of these fundamental difficulties within the context of specific game domains. The interplay between these challenges and the structural properties of different game genres shapes the landscape of Game-MARL research and determines which algorithmic approaches prove most effective in various settings.

3 A Taxonomy of Multi-Agent Game Environments

The landscape of multi-agent game environments has expanded dramatically over the past decade, driven by the increasing demand for challenging, scalable, and reproducible testbeds for reinforcement learning research. A systematic understanding of this landscape is essential for researchers seeking to select appropriate environments for their specific research questions and for interpreting results across different studies. This section provides a comprehensive taxonomy of multi-agent game environments, organizing them along multiple dimensions including game genre, interaction type, information structure, and scale. We examine the defining characteristics of each category, highlight representative environments, and discuss their respective strengths and limitations for multi-agent reinforcement learning research.

3.1 Classification by Game Genre

Game genre provides a natural and intuitive basis for classifying multi-agent environments, as different genres impose distinct structural properties and strategic requirements on agents. The choice of game genre fundamentally shapes the nature of the multi-agent interactions and the algorithmic approaches that are likely to succeed.

3.1.1 Board Games and Classical Games

Board games represent the earliest and most established class of game environments for artificial intelligence research. These games are characterized by discrete state and action spaces, perfect information in many cases, and turn-based mechanics that simplify the temporal dynamics of interaction. The theoretical foundations of multi-agent reinforcement learning trace back to stochastic games formalized by (Shapley, 1953), which were explicitly motivated by the analysis of generalized board games. (Littman, 1994) later operationalized this framework for reinforcement learning, demonstrating how Markov games could be applied to simple grid-world scenarios inspired by board game mechanics.

Chess and Go have served as enduring benchmarks for adversarial multi-agent intelligence. (Silver et al., 2016) achieved a historic breakthrough with AlphaGo, which combined deep neural networks with Monte Carlo tree search to defeat human world champions in Go. This work was extended by (Silver et al., 2017) with AlphaZero, which learned entirely through self-play without any human data, achieving superhuman performance in Go, chess, and shogi. These achievements demonstrated the power of self-play reinforcement learning in zero-sum, perfect-information games, establishing a paradigm that would influence subsequent work in more complex domains.

Beyond the classic two-player zero-sum games, multiplayer board games introduce additional complexity through mixed motives and temporal alliances. The game of Diplomacy, which requires negotiation, cooperation, and betrayal among seven players, has emerged as a challenging testbed for multi-agent reasoning. (Gray et al., 2020) introduced an AI agent for Diplomacy that combined reinforcement learning with human imitations, while (Bakhtin et al., 2022) developed Cicero, which achieved human-level performance by integrating strategic reasoning with natural language negotiation. Hanabi, a cooperative card game where players cannot see their own cards but can see others', has become a standard benchmark for theory of mind and common-sense reasoning in multi-agent systems (Bard et al., 2020).

3.1.2 Real-Time Strategy Games

Real-time strategy games represent a significant leap in complexity compared to turn-based board games. These environments require agents to manage multiple units simultaneously, gather resources, build structures, and engage in tactical combat, all under real-time constraints. The combination of vast state spaces, long time horizons, and mixed cooperative-competitive dynamics makes real-time strategy games particularly challenging and realistic testbeds for multi-agent reinforcement learning.

StarCraft II has emerged as the preeminent benchmark in this category, offering a rich environment for studying hierarchical coordination and multi-agent teamwork. (Vinyals et al., 2019) achieved grandmaster level in *StarCraft II* using a league of reinforcement learning agents, demonstrating that multi-agent training with diverse opponents could produce robust and adaptive strategies. The SMAC benchmark (Samvelyan et al., 2019) extracted micromanagement scenarios from *StarCraft II*, providing a focused testbed for cooperative multi-agent algorithms where a team of units must defeat an opposing team. This benchmark has been widely adopted for evaluating value decomposition methods such as QMIX (Rashid et al., 2018) and its variants.

Other real-time strategy environments have also contributed to the field. *Warcraft III* and its custom map *Defense of the Ancients* laid the groundwork for the multiplayer online battle arena genre, while *StarCraft I* served as an early platform for multi-agent research through the BWAPI interface (Ontañón et al., 2013). More recently, *Neural MMO* (Suarez et al., 2019) introduced a massively multi-agent environment inspired by persistent online worlds, supporting hundreds of agents in a shared virtual space with resource competition and survival objectives.

3.1.3 Multiplayer Online Battle Arena Games

Multiplayer online battle arena games have become one of the most commercially successful and academically challenging genres for multi-agent reinforcement learning. In these games, two teams of typically five players each control powerful characters with unique abilities, competing to destroy the opposing team's base. The combination of real-time action, long-term strategy, team coordination, and partial observability creates a rich multi-agent challenge.

Dota 2 represents the most prominent example in this category. (Berner et al., 2019) developed OpenAI Five, a team of five neural network agents that defeated the world champions in *Dota 2*. This achievement required scaling multi-agent reinforcement learning to handle long time horizons, complex action spaces, and the need for coordinated team play. The OpenAI Five project demonstrated the feasibility of end-to-end learning in complex team games, though it required massive computational resources and careful reward shaping.

Honor of Kings, a mobile multiplayer online battle arena game developed by Tencent, has also served as a platform for industrial-scale multi-agent reinforcement learning. (Ye et al., 2020) developed the JueWu system, which achieved human-expert level performance through a combination of reinforcement learning, imitation learning, and self-play. The system incorporated hierarchical training strategies to manage the complexity of five-on-five team battles with dozens of characters and abilities. Similarly, (Wu et al., 2021) introduced a multi-agent reinforcement learning framework for team-based multiplayer online battle arena games that explicitly modeled inter-agent coordination through attention mechanisms.

3.1.4 First-Person Shooter Games

First-person shooter games immerse agents in three-dimensional environments where they must navigate, aim, and engage opponents in real-time combat. These environments test low-level motor control, spatial awareness, and rapid decision-making alongside high-level strategic reasoning. The partial observability inherent in first-person perspective creates additional challenges for agent coordination.

ViZDoom (Kempka et al., 2016) provides a lightweight interface to the classic game *Doom*, enabling reinforcement learning research in first-person shooter settings. The platform supports both single-agent and multi-agent scenarios, including deathmatch modes where agents compete against each other. (Dosovitskiy and Koltun, 2017) used *ViZDoom* to study learning from raw pixels in adversarial settings, while (Jaderberg et al., 2019) incorporated first-person shooter elements into the broader DeepMind Lab environment.

More recently, *Counter-Strike: Global Offensive* and its successor *Counter-Strike 2* have attracted research interest due to their complex team dynamics, economic systems, and round-based structure. (Pearce et al., 2023) introduced a benchmark for cooperative-competitive multi-agent learning based on *Counter-Strike*, focusing on the coordination challenges inherent in bomb-planting and hostage-rescue scenarios. (Smith et al., 2024) extended this work by studying how agents can learn to communicate through in-game actions and map pings.

3.1.5 Card Games

Card games present unique challenges for multi-agent reinforcement learning due to their inherent partial observability, stochastic elements, and complex information structures. Players must reason about hidden information, infer opponents' holdings, and balance immediate gains against long-term strategic considerations.

Poker has served as a canonical benchmark for imperfect-information games. (Brown and Sandholm, 2019) developed Pluribus, a superhuman AI for six-player no-limit Texas hold'em, demonstrating that self-play reinforcement learning could achieve equilibrium strategies in large-scale imperfect-information games. The success of Pluribus highlighted the importance of search algorithms that operate in information-space rather than state-space.

Hearthstone and *Magic: The Gathering* introduce additional complexity through large card pools, deck-building mechanics, and turn-based but interactive gameplay. (Zhang et al., 2022) studied multi-agent learning in *Hearthstone* battlegrounds mode, where eight players compete through sequential drafting and combat phases. (Chen et al., 2023) developed a reinforcement learning framework for collectible card games that incorporated opponent modeling and deck adaptation.

3.1.6 Open-World Sandbox Games

Open-world sandbox games represent the frontier of multi-agent game environments, offering virtually unlimited possibilities for exploration, interaction, and emergent behavior. These environments test agents' ability to learn diverse skills, set their

own goals, and adapt to novel situations.

Minecraft has emerged as the preeminent platform for open-world artificial intelligence research. The MineRL project (Guss et al., 2019) introduced a large-scale dataset of human demonstrations and a competition focused on sample-efficient reinforcement learning. The Project Malmö platform (Johnson et al., 2016) provided a programmable interface to *Minecraft*, enabling researchers to design custom multi-agent scenarios ranging from cooperative building to competitive survival.

Recent advances have integrated large language models with *Minecraft* agents to achieve unprecedented capabilities. (Wang et al., 2023) introduced Voyager, an open-ended embodied agent that continuously explores *Minecraft* worlds, discovers new skills, and maintains a growing skill library through iterative prompting of large language models. (Fan et al., 2024) extended this approach to multi-agent settings, studying how language-guided agents can coordinate to build complex structures and accomplish tasks that require teamwork.

3.2 Classification by Interaction Type

Beyond game genre, the nature of interactions among agents provides a fundamental axis for classifying multi-agent environments. The interaction type determines the alignment of agent objectives and shapes the appropriate solution concepts and evaluation metrics.

3.2.1 Fully Cooperative Environments

In fully cooperative environments, all agents share a common reward function and work together to achieve shared objectives. These settings test agents' ability to coordinate, communicate implicitly or explicitly, and allocate subtasks effectively. The theoretical framework for cooperative multi-agent reinforcement learning draws on team theory and cooperative game theory, with solution concepts focusing on team-optimal joint policies.

The SMAC benchmark (Samvelyan et al., 2019) provides a canonical example of fully cooperative environments, where a team of *StarCraft II* units must defeat an opposing team controlled by the game's built-in artificial intelligence. Success requires agents to learn emergent coordination behaviors such as focusing fire, flanking maneuvers, and ability combos. The cooperative nature of SMAC has made it the primary testbed for value decomposition methods including VDN (Sunehag et al., 2017), QMIX (Rashid et al., 2018), and QTRAN (Son et al., 2019).

Cooperative card games like Hanabi (Bard et al., 2020) introduce additional challenges through partial observability and the need for theory of mind. In Hanabi, players cannot see their own cards but can see others', requiring agents to reason about what other agents know and to communicate implicitly through the cards they play. This environment has spurred research into deep theory of mind and common-sense reasoning in multi-agent systems.

3.2.2 Fully Competitive Environments

Fully competitive environments pit agents against each other in zero-sum settings where one agent's gain is another's loss. These environments test agents' ability to anticipate opponent moves, exploit weaknesses, and achieve equilibrium strategies. Game theory provides a rich foundation for competitive multi-agent learning, with Nash equilibrium serving as the standard solution concept.

Two-player zero-sum games including chess, Go, and poker represent the classical instances of competitive environments. (Silver et al., 2017) demonstrated that self-play reinforcement learning could achieve superhuman performance in chess and Go without any human data, while (Brown and Sandholm, 2019) extended these techniques to imperfect-information poker games. The success of self-play in competitive environments relies on the fact that agents can generate progressively stronger opponents through their own learning, creating an automatic curriculum.

3.2.3 Mixed-Motive Environments

Mixed-motive environments capture the complexity of real-world multi-agent interactions where incentives are partially aligned and partially conflicting. These environments test agents' ability to form temporary alliances, negotiate, and reason about the strategic intentions of others. The solution concepts become more complex, encompassing correlated equilibria, social welfare optima, and bargaining solutions.

The Melting Pot platform (Agapiou et al., 2022) was explicitly designed to study mixed-motive interactions through a suite of multi-agent scenarios grounded in common-pool resource dilemmas, public goods games, and intergroup conflict. These environments capture fundamental social challenges including cooperation under risk of exploitation, punishment of defectors, and coordination on mutually beneficial outcomes.

(Leibo et al., 2017) studied the emergence of cooperation and competition in mixed-motive environments through the lens of multi-agent reinforcement learning, demonstrating that agents could learn to cooperate when the game structure rewarded reciprocal altruism. This line of research connects multi-agent reinforcement learning to evolutionary game theory and behavioral economics, providing insights into the conditions under which cooperation emerges in populations of learning agents.

3.3 Classification by Information Structure

The information available to agents fundamentally shapes the learning problem and the algorithmic approaches required for effective multi-agent reinforcement learning.

3.3.1 Perfect Information Environments

In perfect information environments, all agents observe the complete state of the game without any hidden information. Chess and Go exemplify perfect information games, as both players can see the entire board state. These environments simplify the learning problem by eliminating the need for belief state maintenance and opponent modeling for information inference.

Perfect information environments have been particularly amenable to search-based methods combined with reinforcement learning. (Silver et al., 2017) leveraged perfect information in AlphaZero to perform Monte Carlo tree search during both training and play, achieving superhuman performance through the combination of learned value functions and forward search.

3.3.2 Imperfect Information Environments

Imperfect information environments introduce hidden information that agents must infer from observations and actions. Poker exemplifies imperfect information, as players cannot see opponents' cards and must infer hand strength from betting patterns. Real-time strategy games introduce fog of war, hiding portions of the map that agents cannot currently observe.

Partially observable stochastic games provide the formal framework for imperfect information environments (Hansen et al., 2004). In these settings, each agent receives observations according to an observation function rather than accessing the full state, leading to the policy representation formalized in Equation (9). The challenge of partial observability has motivated research into recurrent neural networks for maintaining internal state (Hausknecht and Stone, 2015) and into deep theory of mind for reasoning about others' hidden information (Rabinowitz et al., 2018).

3.4 Classification by Scale

The number of agents in an environment introduces distinct challenges and often requires different algorithmic approaches. Environments can be classified by whether they involve small-scale interactions with two to ten agents, or large-scale systems with dozens to hundreds of agents.

Small-scale environments including two-player games and five-player team games dominate the literature due to their tractability and the availability of established benchmarks. These environments allow researchers to focus on algorithmic innovations without the engineering challenges of massive scalability.

Large-scale environments with many agents introduce qualitatively different challenges including the need for scalable coordination mechanisms, spatial and temporal abstraction, and handling of emergent group dynamics. (Suarez et al., 2019) introduced Neural MMO specifically to study large-scale multi-agent learning with hundreds of agents in persistent virtual worlds. (Baker et al., 2019) studied emergent tool use and social dynamics in large populations of agents playing the collaborative game Hide and Seek, demonstrating that complex social behaviors could emerge from simple reinforcement learning objectives.

4 Core Algorithms and Methodologies

The rapid advancement of multi-agent reinforcement learning in game environments has been driven by the development of increasingly sophisticated algorithms that address the fundamental challenges outlined in Section 2.3. This section provides a comprehensive overview of the core algorithmic families that have emerged as dominant approaches in Game-MARL research. We organize these algorithms according to their underlying principles and architectural patterns, examining their theoretical foundations, practical implementations, and relative strengths and weaknesses across different game genres.

4.1 Value Decomposition Methods

Value decomposition methods have emerged as one of the most influential families of algorithms for cooperative multi-agent reinforcement learning. These approaches address the multi-agent credit assignment problem by factoring a global joint action-value function into components associated with individual agents, enabling scalable learning in environments with many agents while maintaining the ability to evaluate joint actions.

4.1.1 Foundations of Value Decomposition

The fundamental insight underlying value decomposition methods is that while agents must coordinate their actions to achieve team objectives, learning can be structured to respect the principle of individual-global max, which requires that the optimal joint action for the team corresponds to the collection of optimal individual actions from each agent's perspective. Formally, this principle states that there exists a monotonic relationship between the joint action-value function and the individual utility functions such that

$$\arg \max_{\mathbf{a}} Q_{\text{tot}}(s, \mathbf{a}) = \left(\arg \max_{a^1} Q_1(s, a^1), \dots, \arg \max_{a^n} Q_n(s, a^n) \right), \quad (11)$$

where Q_{tot} is the joint action-value function and Q_i are individual utility functions for each agent. This principle ensures that agents can select actions greedily with respect to their individual utilities without compromising team optimality, enabling decentralized execution after centralized training.

4.1.2 Value Decomposition Networks

The Value Decomposition Network represents the first deep learning approach to factorization in cooperative multi-agent reinforcement learning (Sunehag et al., 2017). VDN decomposes the joint action-value function as a simple sum of individual agent value functions:

$$Q_{\text{tot}}(s, \mathbf{a}) = \sum_{i=1}^n Q_i(s, a^i), \quad (12)$$

where each Q_i is a neural network that takes the global state s and the agent's individual action a^i as input, or more commonly, takes the agent's local observation o^i and action. The parameters of all individual networks are trained end-to-end to minimize the standard DQN loss using the joint reward signal:

$$\mathcal{L}(\theta) = \mathbb{E}_{(s, \mathbf{a}, r, s') \sim \mathcal{D}} \left[\left(r + \gamma \max_{\mathbf{a}'} Q_{\text{tot}}(s', \mathbf{a}'; \theta^-) - Q_{\text{tot}}(s, \mathbf{a}; \theta) \right)^2 \right], \quad (13)$$

where θ^- represents target network parameters and $\mathcal{D}$ is a replay buffer containing trajectories of joint experiences.

The additive factorization in VDN ensures that the individual-global max principle holds trivially, as the sum of individual maxima equals the maximum of the sum. However, this simple additive structure may be insufficient for problems requiring complex coordination, as it cannot represent interactions where an agent's contribution depends nonlinearly on others' actions.

4.1.3 QMIX and Monotonic Factorization

QMIX addresses the limitations of additive factorization by introducing a more expressive mixing network that combines individual agent utilities into a joint action-value function while enforcing monotonicity constraints (Rashid et al., 2018). The QMIX architecture consists of two components: individual agent networks that generate utilities $Q_i(s, a^i)$, and a mixing network that combines these utilities into $Q_{\text{tot}}(s, \mathbf{a})$.

The mixing network takes the individual utilities as input and produces the joint value through a feedforward network with nonnegative weights:

$$Q_{\text{tot}}(s, \mathbf{a}) = f_{\text{mix}} \left(s, Q_1(s, a^1), \dots, Q_n(s, a^n) \right), \quad (14)$$

subject to the monotonicity constraint:

$$\frac{\partial Q_{\text{tot}}}{\partial Q_i} \geq 0, \quad \forall i \in \mathcal{N}. \quad (15)$$

This constraint is enforced by restricting the weights of the mixing network to be nonnegative, which ensures that the joint action-value function is monotonic in each individual utility. The monotonicity condition is sufficient to guarantee the individual-global max principle while allowing for more complex interactions than simple addition.

The mixing network also incorporates the global state s through hypernetworks that generate the weights and biases of the mixing layers, enabling the coordination pattern to adapt to different environmental contexts. This architectural choice allows QMIX to represent a broader class of cooperative tasks than VDN while maintaining tractable learning.

Empirical evaluations on the SMAC benchmark demonstrate that QMIX significantly outperforms VDN and independent learners across a range of cooperative micromanagement tasks (Samvelyan et al., 2019). The algorithm's success has spawned numerous extensions and variants that address remaining limitations.

4.1.4 Advanced Value Decomposition Methods

Following the success of QMIX, several advanced value decomposition methods have been developed to address specific limitations and extend the representational capacity of factorization approaches.

QTRAN addresses a fundamental limitation of monotonic factorization by recognizing that not all cooperative tasks satisfy the monotonicity condition (Son et al., 2019). QTRAN proposes a transformation that converts the original joint action-value function into an equivalent representation that is additively decomposable while preserving the optimal joint actions. The algorithm maintains three value functions: the original Q_{tot} , an additive decomposition $Q_{\text{add}} = \sum_i Q_i$, and a state-value network $V(s)$ that captures the gap between them. The learning objective enforces that Q_{tot} and Q_{add} coincide for optimal actions while allowing divergence for suboptimal joint actions. This approach enables QTRAN to represent tasks that require non-monotonic coordination, such as the classic two-step matrix game where agents must coordinate on complementary actions.

Weighted QMIX extends the monotonic framework by introducing weighting factors that emphasize transitions where the current joint action is optimal (Rashid et al., 2020). The weighted loss function is given by

$$\mathcal{L}(\theta) = \mathbb{E} \left[w(s, \mathbf{a}) \left(r + \gamma \max_{\mathbf{a}'} Q_{\text{tot}}(s', \mathbf{a}'; \theta^-) - Q_{\text{tot}}(s, \mathbf{a}; \theta) \right)^2 \right], \quad (16)$$

where $w(s, \mathbf{a})$ assigns higher weight to experiences where the taken joint action is closer to optimal. This weighting scheme improves learning efficiency by focusing updates on transitions that provide the most informative feedback.

QPLEX introduces a duplex dueling network architecture that factorizes the joint action-value function into advantage and value components (Wang et al., 2021). The approach maintains the individual-global max principle while enabling full expressiveness through a carefully designed transformation. QPLEX achieves state-of-the-art performance on challenging SMAC scenarios that require complex coordination, including tasks where agents must synchronize their attacks or execute precisely timed maneuvers.

4.1.5 Theoretical Insights and Limitations

The theoretical understanding of value decomposition methods has advanced significantly in recent years. (Mahajan et al., 2021) analyzed the representational power of monotonic factorization, demonstrating that while monotonic networks can represent any function over the joint action space, the constraints imposed during learning may limit practical performance. (Wang et al., 2020) established connections between value decomposition and counterfactual reasoning, showing that individual utilities in QMIX can be interpreted as counterfactual baselines that measure each agent’s contribution to team success.

Despite their success, value decomposition methods face several persistent limitations. The assumption of centralized training with access to global state may be unrealistic in some applications where communication constraints prevent sharing complete information. Additionally, these methods are primarily designed for fully cooperative settings and do not naturally extend to competitive or mixed-motive environments. The scalability of value decomposition to very large numbers of agents also remains challenging, as the mixing network must process all individual utilities simultaneously.

4.2 Centralized Training with Decentralized Execution

Centralized Training with Decentralized Execution has emerged as a dominant paradigm in multi-agent reinforcement learning, addressing the tension between the need for coordinated learning and the requirement for independent execution. Under this paradigm, agents have access to additional information during training that is unavailable during execution, enabling more effective learning while maintaining decentralized policies at test time.

4.2.1 Foundations of CTDE

The CTDE framework recognizes that training and deployment conditions may differ in multi-agent systems. During training, it is often feasible to access global state information, share gradients across agents, and coordinate exploration strategies. However, during execution, agents may be limited to local observations and must act independently without communication. CTDE algorithms exploit this asymmetry by using privileged information during training to learn better decentralized policies.

Formally, CTDE algorithms learn policies $\pi^i(a^i|o^i)$ that depend only on local observations o^i , but during training, they have access to additional information such as the global state s , other agents’ observations, or inter-agent communication channels. This paradigm has proven particularly effective in partially observable environments where local observations provide insufficient information for optimal decision-making.

4.2.2 MADDPG

The Multi-Agent Deep Deterministic Policy Gradient algorithm extends the deep deterministic policy gradient framework to multi-agent settings, providing a general-purpose CTDE algorithm for environments with continuous action spaces and mixed cooperative-competitive interactions (Lowe et al., 2017). MADDPG adopts a centralized critic, decentralized actor architecture where each agent maintains its own policy π_{θ^i} and a centralized critic Q_{ϕ^i} that evaluates actions given the global state and all agents’ actions.

The centralized critic for agent i is trained to minimize the loss:

$$\mathcal{L}(\phi^i) = \mathbb{E}_{\mathcal{D}} \left[\left(Q_{\phi^i}(s, \mathbf{a}) - \left(r^i + \gamma Q_{\phi^i}(s', \mathbf{a}') \right) \Big|_{\mathbf{a}' = \pi_{\theta}(s')} \right)^2 \right], \quad (17)$$

where $\pi_{\theta}(s')$ denotes the vector of target policies evaluated at the next state. The actor for agent i is updated using the policy gradient:

$$\nabla_{\theta^i} J(\theta^i) = \mathbb{E}_{\mathcal{D}} \left[\nabla_{\theta^i} \pi_{\theta^i}(a^i|o^i) \nabla_{a^i} Q_{\phi^i}(s, \mathbf{a}) \Big|_{a^i = \pi_{\theta^i}(o^i)} \right]. \quad (18)$$

The use of centralized critics that condition on all agents’ actions addresses the non-stationarity problem by providing each agent with a stable learning signal that accounts for the policies of other agents. During execution, agents act based only on their local observations and policies, enabling decentralized deployment.

MADDPG has been successfully applied to a variety of multi-agent game environments, including cooperative communication tasks, competitive predator-prey scenarios, and mixed-motive physical coordination games. The algorithm’s flexibility in handling arbitrary reward structures makes it particularly suitable for mixed-motive environments where agents may have partially aligned objectives.

4.2.3 COMA

Counterfactual Multi-Agent Policy Gradients introduces a counterfactual baseline mechanism specifically designed for cooperative multi-agent settings (Foerster et al., 2018). COMA addresses the multi-agent credit assignment problem by computing a baseline that marginalizes out a single agent's action while keeping others' actions fixed, enabling each agent to estimate its individual contribution to the team reward.

The core innovation in COMA is the counterfactual advantage function for agent i :

$$A^i(s, \mathbf{a}) = Q(s, \mathbf{a}) - \sum_{a^{i'}} \pi^i(a^{i'} | o^i) Q(s, (\mathbf{a}^{-i}, a^{i'})), \quad (19)$$

where $\mathbf{a}^{-i}$ denotes the actions of all agents except i . This advantage measures how much better the taken action a^i is compared to the average action agent i could have taken, given that all other agents execute their chosen actions. The baseline is efficiently computed using a single forward pass through a critic network that outputs Q-values for all possible actions of agent i simultaneously.

COMA employs a centralized critic that conditions on the global state and all agents' actions, similar to MADDPG, but the critic architecture is designed to output Q-values for each agent's entire action space in a single forward pass. This design enables efficient computation of the counterfactual baseline and reduces the variance of policy gradient estimates.

Empirical evaluations on the StarCraft micromanagement tasks demonstrate that COMA significantly outperforms independent policy gradient methods and achieves performance competitive with value decomposition approaches (Samvelyan et al., 2019). The counterfactual baseline mechanism provides effective credit assignment even in tasks with delayed rewards and complex coordination requirements.

4.2.4 MAPPO

Multi-Agent Proximal Policy Optimization adapts the highly successful PPO algorithm to multi-agent settings, combining the stability and sample efficiency of PPO with the CTDE paradigm (Yu et al., 2022). MAPPO maintains a centralized value function that conditions on global state information while learning decentralized policies that depend only on local observations.

The MAPPO objective for agent i extends the standard PPO clipped surrogate objective to the multi-agent setting:

$$\mathcal{L}_{\text{CLIP}}(\theta^i) = \mathbb{E} \left[\min \left(r_t(\theta^i) \hat{A}_t^i, \text{clip}(r_t(\theta^i), 1 - \epsilon, 1 + \epsilon) \hat{A}_t^i \right) \right], \quad (20)$$

where $r_t(\theta^i) = \frac{\pi_{\theta^i}(a_t^i | o_t^i)}{\pi_{\theta^i}^{\text{old}}(a_t^i | o_t^i)}$ is the probability ratio, and $\hat{A}_t^i$ is an advantage estimate computed using the centralized value function. The centralized value function $V_\phi(s_t)$ is trained to minimize the squared error with respect to empirical returns:

$$\mathcal{L}(\phi) = \mathbb{E} \left[(V_\phi(s_t) - \hat{R}_t)^2 \right], \quad (21)$$

where $\hat{R}_t$ is the discounted return from time t .

A key insight from (Yu et al., 2022) is that many of the design choices that make PPO successful in single-agent settings, such as value function normalization, reward scaling, and proper handling of recurrent neural networks, are equally important in multi-agent settings. MAPPO incorporates these elements and achieves strong performance across a range of cooperative multi-agent benchmarks, often matching or exceeding the performance of more complex algorithms.

The simplicity and stability of MAPPO have made it a popular choice for multi-agent reinforcement learning research. The algorithm scales well to many agents and has been successfully applied to challenging game environments including SMAC and Google Research Football.

4.2.5 Comparison of CTDE Methods

The three CTDE methods discussed above represent distinct approaches to the multi-agent learning problem. MADDPG offers the greatest flexibility in terms of reward structure, handling cooperative, competitive, and mixed-motive settings with equal ease. However, its reliance on deterministic policies and experience replay can lead to instability in some environments. COMA provides theoretically grounded credit assignment for cooperative tasks but requires a critic architecture that scales with the action space size, limiting its applicability to domains with large discrete action spaces. MAPPO combines the stability of PPO with the scalability of decentralized policies, making it an attractive choice for many applications, though its performance depends critically on careful hyperparameter tuning.

Recent work has begun to unify insights across these approaches. (Lyu et al., 2021) proposed a framework for understanding CTDE algorithms through the lens of variance reduction, showing that the success of these methods stems primarily from their ability to reduce gradient variance through centralized value functions. (Kuba et al., 2022) introduced a unified view of multi-agent policy gradients that encompasses MADDPG, COMA, and MAPPO as special cases, providing theoretical insights into their relative strengths and weaknesses.

4.3 Communication-Based Methods

In many multi-agent game environments, agents can benefit from explicit communication to coordinate their actions, share information about partially observed states, and negotiate strategies. Communication-based methods extend the multi-agent

reinforcement learning framework by allowing agents to exchange messages, either through predefined communication channels or through emergent communication protocols learned during training.

4.3.1 Learning to Communicate

The challenge of learning communication protocols alongside action policies has attracted significant research interest. Early work by (Sukhbaatar et al., 2016) introduced CommNet, a neural architecture that enables agents to communicate through a continuous vector space. At each time step, agents compute a hidden state based on their local observation and messages received from other agents:

$$h_t^i = f \left(o_t^i, \frac{1}{n-1} \sum_{j \neq i} h_{t-1}^j \right), \quad (22)$$

where f is a neural network, typically a gated recurrent unit. The messages are simply the hidden states themselves, enabling differentiable communication channels that can be trained end-to-end with backpropagation through time.

CommNet demonstrated that agents could learn to coordinate in tasks requiring information sharing, such as the Traffic Junction and predator-prey domains. However, the continuous nature of the communication channel and the homogeneous architecture limit the interpretability of the learned messages.

4.3.2 Structured Communication Protocols

Subsequent work has introduced more structured communication protocols that address the limitations of early approaches. TarMAC extends the communication framework by incorporating attention mechanisms that enable agents to selectively attend to messages from other agents based on their relevance (Das et al., 2019). The TarMAC architecture computes messages as:

$$m_t^i = \text{MLP} \left(h_t^i \right), \quad (23)$$

$$\alpha_t^{i \rightarrow j} = \text{softmax} \left((W_q h_t^i)^T (W_k m_t^j) \right), \quad (24)$$

$$h_t^{i, \text{new}} = \text{GRU} \left(h_t^i, \sum_{j \neq i} \alpha_t^{i \rightarrow j} W_v m_t^j \right), \quad (25)$$

where W_q , W_k , and W_v are learned projection matrices. The attention mechanism allows agents to focus on messages from relevant peers, improving scalability and coordination in large multi-agent systems.

IC3Net introduces gating mechanisms that enable agents to learn when to communicate, addressing the problem of communication overhead and enabling more efficient information exchange (Singh et al., 2019). Each agent maintains a binary communication gate:

$$g_t^i = \sigma \left(\text{MLP} \left(h_t^i \right) \right), \quad (26)$$

and the message sent by agent i is $g_t^i \cdot m_t^i$. This gating mechanism allows agents to develop selective communication strategies, only broadcasting information when it is likely to be useful to others.

4.3.3 Emergent Communication and Language

A growing body of research investigates the emergence of communication protocols that resemble natural language. These approaches typically place agents in environments where communication is necessary for task completion and evaluate the properties of the resulting protocols.

(Lazaridou et al., 2017) studied emergent communication in a referential game where a sender agent must communicate a target image to a receiver agent. The agents develop communication protocols that exhibit compositionality and systematicity, properties that are characteristic of natural language. (Evtimova et al., 2018) extended this work to multi-step communication, showing that agents could learn to ask clarifying questions and provide incremental information.

The challenge of grounding emergent communication in shared experience has been addressed through various mechanisms. (Mordatch and Abbeel, 2018) proposed a framework for emergent language learning that combines reinforcement learning with self-supervision, enabling agents to develop communication protocols that are grounded in their shared environment. (Lowe et al., 2020) investigated the factors that promote the emergence of compositional communication, finding that population size, environment complexity, and the presence of common ground all influence the properties of learned protocols.

Recent work has begun to integrate large language models into multi-agent communication frameworks. (Park et al., 2023) introduced a framework where agents use natural language generated by large language models to coordinate in complex game environments, enabling human-interpretable communication and flexible coordination strategies. (Zhang et al., 2024) extended this approach to enable agents to negotiate and form alliances through natural language dialogue.

4.3.4 Challenges in Communication-Based Methods

Despite significant progress, communication-based methods face several persistent challenges. The scalability of communication to many agents remains problematic, as fully connected communication graphs become prohibitive for large populations. Sparse communication mechanisms and hierarchical communication structures offer potential solutions but require further development.

The interpretability of learned communication protocols remains limited, particularly when messages are represented as continuous vectors. Methods that encourage discrete communication or that ground communication in shared representations may improve interpretability but often at the cost of reduced flexibility.

The trade-off between communication and autonomy presents another challenge. Agents that communicate frequently may achieve better coordination but incur communication costs and may become overly dependent on communication channels that could fail during deployment. Learning when to communicate and what information to share remains an active area of research.

4.4 League Training and Population-Based Methods

In competitive and mixed-motive game environments, learning against a fixed set of opponents can lead to overfitting and brittle strategies that exploit specific weaknesses rather than learning robust, generalizable behaviors. League training and population-based methods address this challenge by maintaining diverse populations of agents and orchestrating matches between them to generate an automatic curriculum of increasingly challenging opponents.

4.4.1 Self-Play and Its Limitations

Self-play, where an agent trains by playing against copies of itself, has proven remarkably effective in two-player zero-sum games. (Silver et al., 2017) demonstrated that self-play alone could produce superhuman performance in chess and Go, with AlphaZero learning entirely through games against itself. The theoretical foundation for self-play lies in the fact that in zero-sum games, playing against increasingly strong versions of oneself generates a natural curriculum toward optimal play.

However, self-play has significant limitations in more complex multi-agent settings. In games with more than two players or with non-zero-sum payoffs, self-play may converge to cyclic behaviors or fail to explore the full space of strategies. Additionally, self-play can lead to specialization on a narrow set of strategies, leaving agents vulnerable to out-of-distribution opponents encountered during deployment.

4.4.2 AlphaStar League Training

The AlphaStar system for StarCraft II introduced a sophisticated league training framework that addressed the limitations of simple self-play (Vinyals et al., 2019). The league consisted of three types of agents: main agents that were trained to maximize performance against the full league, league explorers whose objective was to discover weaknesses in the main agents, and main past agents that were snapshot copies of previous main agents.

The training process orchestrated matches between all agents in the league, with each agent updating its policy based on its match outcomes. The diverse population ensured that main agents encountered a wide variety of strategies, promoting robust and generalizable play. The league explorers specifically targeted weaknesses in main agents, creating an automatic curriculum that continuously challenged the main agents to improve.

The success of league training in AlphaStar demonstrated the power of population-based methods for complex multi-agent games. The resulting agent achieved grandmaster level in StarCraft II, defeating professional human players through a combination of strategic diversity and tactical precision.

4.4.3 OpenAI Five and Prioritized Training

OpenAI Five, the system that achieved superhuman performance in Dota 2, employed a related but distinct approach to population-based training (Berner et al., 2019). The system maintained a fixed pool of agent versions and used a prioritized sampling mechanism to select opponents during training.

The prioritization mechanism assigned higher probability to selecting opponents that provided informative learning signals, such as those that recently defeated the current agent or those that exhibited complementary strengths. This dynamic opponent selection focused training on the most challenging and informative matchups, improving sample efficiency and promoting robust learning.

OpenAI Five also employed a technique called surgery to scale the league to distributed training across thousands of machines. The system periodically paused training to evaluate all agent versions against each other, updating a win-rate matrix that informed the prioritization mechanism for subsequent training iterations.

4.4.4 PSRO and Game-Theoretic Methods

Policy Space Response Oracles provides a game-theoretic framework for population-based multi-agent learning (Lanctot et al., 2017). PSRO maintains a population of policies and iteratively computes best responses to the current population, expanding the population with new policies that exploit weaknesses in existing ones.

The PSRO algorithm proceeds as follows: initialize a population $\Pi = \{\pi_0\}$; at each iteration, compute a best response π_{BR} to the meta-Nash equilibrium over Π ; add π_{BR} to Π ; update the meta-Nash equilibrium over the expanded population. This process continues until convergence, producing a population of policies that approximates a Nash equilibrium of the full game.

PSRO has been extended in numerous directions. (Muller et al., 2020) introduced rectified PSRO, which focuses best response computation on strategies that are likely to improve upon existing ones. (McAleer et al., 2021) developed pipeline PSRO, which enables scalable training by parallelizing best response computation. (Marris et al., 2022) unified PSRO with league training through a common framework based on empirical game-theoretic analysis.

4.4.5 Population-Based Training

Population-Based Training offers a complementary approach to multi-agent learning that focuses on optimizing hyperparameters and learning dynamics alongside policy parameters (Jaderberg et al., 2017). In PBT, a population of agents trains in parallel, with each agent periodically copying parameters from better-performing agents and exploring hyperparameter variations through random perturbations.

While originally developed for single-agent settings, PBT has been successfully applied to multi-agent domains. The approach enables automatic adaptation of learning rates, exploration parameters, and network architectures, reducing the burden of manual hyperparameter tuning and often leading to better final performance.

4.5 Algorithms for Mixed-Motive Environments

Mixed-motive environments, where agents have partially aligned and partially conflicting objectives, present unique challenges that require algorithmic approaches beyond those designed for purely cooperative or purely competitive settings. These environments are particularly relevant to many modern multiplayer games and real-world multi-agent systems.

4.5.1 Social Dilemmas and Reciprocity

Social dilemmas, situations where individual rationality leads to collectively suboptimal outcomes, provide a canonical testbed for mixed-motive multi-agent learning. (Leibo et al., 2017) introduced sequential social dilemmas, extending classic matrix-game dilemmas to temporally extended environments where agents interact repeatedly.

In sequential social dilemmas, agents must learn strategies that balance immediate individual gains against long-term collective welfare. (Peysakhovich and Lerer, 2018) demonstrated that agents could learn to cooperate in social dilemmas through appropriate reward shaping and the use of reciprocity mechanisms. The key insight was that agents must be able to recognize defection and respond in ways that discourage future defection, effectively learning tit-for-tat-like strategies.

4.5.2 Opponent Modeling and Intention Recognition

In mixed-motive environments, understanding and predicting opponent behavior becomes essential for effective decision-making. Opponent modeling methods aim to learn models of other agents' policies, enabling more informed action selection and strategic adaptation.

(Raileanu et al., 2018) introduced a framework for opponent modeling in multi-agent reinforcement learning where agents maintain explicit models of other agents' policies and use these models to simulate potential futures. The approach improves performance in competitive and mixed-motive domains by enabling agents to anticipate opponent moves and select counter-strategies accordingly.

(Grover et al., 2018) developed a best-response learning algorithm that alternates between updating agent policies and updating opponent models, ensuring that agents remain robust to changes in opponent behavior. The algorithm achieves strong performance in iterated games and mixed-motive environments where opponents may adapt over time.

4.5.3 Social Value Alignment

Recent work has begun to address the challenge of aligning agent behavior with human values and social norms in mixed-motive environments. (McKee et al., 2020) introduced the concept of social diversity in multi-agent learning, showing that populations with diverse reward structures can achieve better collective outcomes than homogeneous populations.

(Hughes et al., 2018) investigated inequity aversion in multi-agent systems, demonstrating that agents equipped with inequity-averse reward functions learn to achieve fairer outcomes in social dilemmas. This line of research connects multi-agent reinforcement learning to behavioral economics and social psychology, providing insights into how artificial agents can learn to behave in socially desirable ways.

4.6 Algorithmic Frameworks and Design Patterns

Across the diverse landscape of multi-agent reinforcement learning algorithms, several recurring design patterns have emerged that provide useful abstractions for understanding and developing new approaches.

4.6.1 Actor-Critic Architectures

Actor-critic methods, which maintain separate policy and value function representations, have proven particularly effective in multi-agent settings. The separation enables centralized critics that can leverage global information during training while maintaining decentralized actors for execution, as exemplified by MADDPG, COMA, and MAPPO.

The actor-critic pattern provides natural solutions to the multi-agent credit assignment problem by using value functions to evaluate joint actions and attribute credit to individual agents. Counterfactual baselines, as in COMA, represent a particularly sophisticated instantiation of this pattern.

4.6.2 Attention and Transformer Mechanisms

Attention mechanisms have emerged as a powerful tool for managing interactions in multi-agent systems. The ability to selectively attend to relevant agents enables scalable coordination in large populations and adaptive responses to changing interaction patterns.

Transformer architectures, which generalize attention to handle sequences of inputs, have begun to appear in multi-agent reinforcement learning. (Wen et al., 2022) introduced a transformer-based architecture for multi-agent value decomposition that enables flexible handling of variable numbers of agents and adaptive coordination patterns.

4.6.3 Hierarchical and Temporal Abstraction

Hierarchical approaches, which learn policies at multiple levels of temporal abstraction, address the challenge of long time horizons in complex game environments. By learning skills and options that persist over many time steps, hierarchical methods enable more efficient exploration and more effective credit assignment.

(Tang et al., 2018) introduced hierarchical multi-agent reinforcement learning with option frameworks, demonstrating improved performance in tasks requiring extended coordination. (Kim et al., 2020) extended these ideas to enable agents to learn communication protocols at multiple temporal scales, with high-level messages conveying strategic intentions and low-level messages coordinating immediate actions.

The diversity of algorithmic approaches to Game-MARL reflects the richness and complexity of multi-agent game environments. Value decomposition methods provide principled solutions to cooperative credit assignment, CTDE algorithms offer flexible frameworks for mixed cooperative-competitive settings, communication-based methods enable explicit information sharing, and population-based approaches generate robust strategies through diverse training. Understanding the strengths and limitations of each algorithmic family is essential for selecting appropriate methods for specific game domains and research questions.

5 Key Applications and Milestones

The theoretical advances in multi-agent reinforcement learning algorithms have been accompanied by a series of landmark applications that have progressively pushed the boundaries of what is achievable in complex game environments. These milestones not only demonstrate the practical capabilities of Game-MARL but also reveal fundamental insights about learning dynamics, strategic reasoning, and the nature of multi-agent intelligence. This section provides a comprehensive examination of the most significant application domains, tracing the evolution from simple board games to complex open-world environments.

5.1 Classic Board Games: Foundations of Multi-Agent Reasoning

Board games have served as the proving ground for artificial intelligence since the earliest days of the field, offering well-defined rules, perfect information in many cases, and clear metrics for success. The progression of achievements in board games chronicles the development of increasingly sophisticated multi-agent reasoning capabilities.

5.1.1 Chess: Deep Blue and Beyond

The defeat of world champion Garry Kasparov by IBM's Deep Blue in 1997 marked the first major milestone in game-playing artificial intelligence (Campbell et al., 2002). Deep Blue combined massive parallel search with handcrafted evaluation functions and opening books, demonstrating that brute-force computation coupled with domain expertise could achieve superhuman performance in a game of immense complexity. While Deep Blue did not employ reinforcement learning in the modern sense, it established the importance of search and evaluation in adversarial multi-agent settings.

The subsequent development of chess engines followed a trajectory of increasing sophistication, with programs like Stockfish incorporating advanced pruning techniques and endgame tablebases. However, these systems remained fundamentally dependent on human knowledge encoded in evaluation functions and search heuristics.

5.1.2 Go: AlphaGo and AlphaZero

The game of Go presented a fundamentally greater challenge than chess due to its larger branching factor and the difficulty of position evaluation. For decades, the computer Go community viewed superhuman performance as a distant goal requiring fundamental breakthroughs in artificial intelligence.

AlphaGo achieved this breakthrough in 2016, defeating world champion Lee Sedol in a five-game match (Silver et al., 2016). The system combined deep neural networks with Monte Carlo tree search in a novel architecture that learned from both human expert games and self-play. The policy network reduced the effective branching factor by selecting promising moves, while the value network provided position evaluation without exhaustive search. The famous "move 37" in game two, a creative play that surprised expert commentators, demonstrated that the system had discovered novel strategies beyond the human corpus.

AlphaZero extended this approach by eliminating the need for human data entirely (Silver et al., 2017). Through pure self-play reinforcement learning, AlphaZero achieved superhuman performance in Go, chess, and shogi within 24 hours of training. The system's success demonstrated the power of self-play as a general-purpose mechanism for generating training data and the effectiveness of the combination of neural networks with search. The resulting playing style in chess differed markedly from traditional engines, prioritizing long-term positional compensation over immediate material gain.

The AlphaZero architecture introduced several key innovations for multi-agent learning. The use of self-play ensured that agents encountered progressively stronger opponents, creating an automatic curriculum. The integration of search with learned value functions enabled efficient exploration of the game tree while maintaining computational tractability. The representational power of deep neural networks allowed the system to capture subtle patterns and positional concepts that had previously required human expertise.

5.1.3 Multi-Player Board Games: Diplomacy and Hanabi

While two-player zero-sum games yield to self-play and search, multiplayer board games introduce additional complexities including alliance formation, negotiation, and theory of mind. Two games have emerged as particularly significant benchmarks for these challenges.

Diplomacy requires seven players to negotiate alliances, coordinate attacks, and potentially betray partners, all while managing a complex board position. The game has no hidden information but immense strategic complexity arising from the need to reason about other players' intentions and commitments. (Gray et al., 2020) developed an artificial intelligence agent that combined reinforcement learning with imitation learning from human games, achieving strong performance in online play. However, the agent's inability to engage in natural language negotiation limited its strategic sophistication.

Cicero represented a breakthrough in this domain, achieving human-level performance in Diplomacy by combining strategic reasoning with natural language generation (Bakhtin et al., 2022). The system maintained a dialogue model that could generate and interpret messages, enabling it to negotiate alliances, share information, and coordinate plans with human players. Cicero's architecture integrated several components: a planning module that computed strategic values for board positions, a dialogue module that generated messages consistent with the agent's intentions, and a theory-of-mind module that modeled other players' beliefs and intentions. The system's success demonstrated that artificial intelligence could engage in the subtle social reasoning required for multiplayer strategic interaction.

Hanabi presents a different challenge: cooperative play under partial observability. Players cannot see their own cards but can see others', requiring them to reason about what other players know and to communicate implicitly through the cards they play. (Bard et al., 2020) introduced Hanabi as a challenge problem for artificial intelligence, highlighting the need for theory of mind and common-sense reasoning. Subsequent work has developed agents capable of human-level cooperation through various mechanisms including value decomposition, Bayesian inference, and emergent communication protocols.

5.2 Real-Time Strategy: StarCraft and the Challenge of Macro-Strategy

Real-time strategy games represent a quantum leap in complexity compared to turn-based board games. The combination of real-time constraints, vast state spaces, long time horizons, and mixed cooperative-competitive dynamics creates a uniquely challenging testbed for multi-agent reinforcement learning.

5.2.1 The StarCraft AI Competition

StarCraft: Brood War served as the primary platform for real-time strategy artificial intelligence research for over a decade through the Brood War Application Programming Interface (Ontañón et al., 2013). The annual AIIDE competition attracted participants from around the world, driving incremental progress in scripted agents, search-based methods, and early reinforcement learning approaches.

The competition revealed fundamental challenges that would shape subsequent research. Successful agents required robust build orders, effective micromanagement of individual units, strategic adaptation to opponent actions, and economical resource management. Scripted agents with carefully tuned heuristics dominated the competition for years, while learning-based approaches struggled with the sample complexity and exploration requirements of the game.

5.2.2 AlphaStar: Grandmaster Level in StarCraft II

AlphaStar achieved grandmaster level in StarCraft II in 2019, representing a landmark achievement in multi-agent reinforcement learning (Vinyals et al., 2019). The system's architecture incorporated numerous innovations that addressed the unique challenges of real-time strategy games.

The observation space in StarCraft II presents enormous complexity: the agent must process hundreds of units, buildings, and terrain features, each with multiple attributes and capabilities. AlphaStar employed a deep neural network with transformer architectures and pointer networks to handle variable-length inputs and outputs. The network processed spatial features through convolutional layers, unit attributes through attention mechanisms, and produced action distributions over a hierarchical action space.

The action space itself required careful design. StarCraft II actions include camera movements, unit selections, ability usage, and production orders, creating a combinatorial explosion of possibilities. AlphaStar addressed this through a hierarchical action representation with separate heads for action type, selected units, target locations, and other parameters. The network could select actions at a frequency of 22.4 cycles per second, matching human capabilities.

Training AlphaStar required massive computational resources and sophisticated multi-agent learning techniques. The system employed a league training framework with three agent types: main agents trained to maximize performance, league explorers tasked with discovering weaknesses, and past checkpoints providing a diverse set of opponents. This population-based approach generated an automatic curriculum that continuously challenged agents to improve.

The league consisted of up to 900 simultaneously training agents, generating millions of games and years of experience per day. The training process required careful balancing to maintain diversity and prevent convergence to narrow strategies.

Evaluation matches between agents determined the composition of the league and the prioritization of training resources.

AlphaStar's final performance demonstrated superhuman capability across all three StarCraft II races, achieving Grandmaster rank on the European ladder. Analysis of the agent's play revealed sophisticated strategies including multi-pronged attacks, temporal distraction, and adaptive responses to opponent builds. The system's success validated the scalability of multi-agent reinforcement learning to complex real-time domains and established league training as a paradigm for generating robust strategies.

5.2.3 SMAC: A Focused Benchmark for Cooperative Learning

While AlphaStar addressed the full complexity of StarCraft II, the SMAC benchmark extracted focused micromanagement scenarios for systematic evaluation of cooperative multi-agent algorithms (Samvelyan et al., 2019). In SMAC, a team of allied units must defeat an opposing team controlled by the game's built-in artificial intelligence, with scenarios ranging from simple symmetric matchups to complex heterogeneous battles.

The SMAC benchmark has become the standard testbed for value decomposition methods and cooperative multi-agent reinforcement learning. Each scenario specifies unit compositions, map layouts, and victory conditions, enabling controlled comparisons between algorithms. The benchmark includes scenarios requiring diverse coordination patterns: focus fire on high-value targets, ability combos that require precise timing, positioning to exploit terrain advantages, and kiting behaviors that maintain distance from melee units.

Extensive evaluations on SMAC have revealed the relative strengths of different algorithmic approaches. QMIX and its variants consistently outperform independent learners and simple value decomposition methods, particularly in scenarios requiring complex coordination (Rashid et al., 2018; Wang et al., 2021). MAPPO achieves competitive performance with simpler implementation and greater stability (Yu et al., 2022). The benchmark continues to drive algorithmic innovation, with recent work addressing limitations in sample efficiency, exploration, and generalization to unseen scenarios.

5.3 Multiplayer Online Battle Arena: Dota and Team Coordination

Multiplayer online battle arena games amplify the coordination challenges of real-time strategy by requiring five-player teams to work together against opposing teams. The genre combines individual mechanical skill with team-level strategy, creating rich opportunities for multi-agent learning.

5.3.1 OpenAI Five: Defeating World Champions in Dota 2

OpenAI Five achieved superhuman performance in Dota 2 in 2019, defeating the world champions OG in a live exhibition match (Berner et al., 2019). The system represented a team of five neural network agents trained entirely through self-play reinforcement learning.

The technical challenges in Dota 2 are staggering. The game features over 100 heroes with unique abilities, thousands of possible item combinations, and matches lasting up to an hour with continuous real-time action. The observation space includes thousands of entity states, and the action space encompasses millions of possibilities per second. OpenAI Five addressed these challenges through careful system design and massive scaling.

Each agent in OpenAI Five was a long short-term memory network processing a curated observation space of 20,000 numbers representing game state. The network output action probabilities through a hierarchical action head that first selected the type of action and then its parameters. Agents shared weights but maintained separate long short-term memory states, enabling role specialization while leveraging shared representations.

Training proceeded through a distributed reinforcement learning system running on thousands of machines. The system generated 180 years of gameplay experience per day through parallel self-play matches. A novel technique called surgery enabled dynamic adjustment of the training curriculum by resetting agent versions that underperformed against the current population.

The final model demonstrated sophisticated team coordination including smoke ganks where multiple agents become invisible to ambush opponents, Roshan contests requiring precise timing and positioning, and buyback strategies that trade immediate gold for tactical advantage. The system's success validated the scalability of multi-agent reinforcement learning to the highest levels of team-based competition.

5.3.2 JueWu: Industrial-Scale Multi-Agent Learning

Tencent's JueWu system achieved expert-level performance in Honor of Kings, a mobile multiplayer online battle arena game with hundreds of millions of players (Ye et al., 2020). The system integrated reinforcement learning with imitation learning from human expert data, achieving superhuman performance while maintaining human-like playing style.

JueWu's architecture incorporated several innovations for industrial-scale deployment. A hierarchical training framework decomposed the game into macro-strategy and micro-execution components, enabling more efficient learning. The system employed a multi-agent distributed training platform supporting thousands of simultaneous agents with sophisticated resource management.

The resulting agent demonstrated capabilities beyond simple mechanical skill, including strategic drafting of team compositions, adaptive responses to opponent strategies, and coordinated team fights with precise ability timing. JueWu's success demonstrated the feasibility of deploying multi-agent reinforcement learning in commercial game environments with real-world impact.

5.3.3 Coordination Mechanisms in MOBA Games

Research on multiplayer online battle arena games has illuminated fundamental mechanisms for multi-agent coordination. (Wu et al., 2021) introduced attention-based coordination mechanisms that enable agents to focus on relevant teammates during team fights, improving synergy and reducing redundant actions. The attention weights learned by the system revealed emergent role specialization, with support heroes attending to carry heroes while damage dealers focused on enemy positions.

The challenge of credit assignment in multiplayer online battle arena games is particularly acute due to delayed rewards and complex causal chains. A kill five minutes into the game may result from a series of decisions including lane assignments, item purchases, and vision control. (Yang et al., 2022) developed counterfactual credit assignment methods tailored to multiplayer online battle arena games, enabling more efficient learning from sparse team rewards.

5.4 First-Person Shooters: Real-Time Coordination and Spatial Reasoning

First-person shooter games test agents' ability to navigate three-dimensional environments, maintain spatial awareness, and coordinate in real-time combat. These environments bridge the gap between abstract strategic reasoning and embodied interaction.

5.4.1 ViZDoom and Deep Learning from Pixels

ViZDoom provided an accessible platform for reinforcement learning research in first-person shooter environments (Kempka et al., 2016). The environment's lightweight implementation and customizable scenarios enabled systematic study of learning from raw pixels in adversarial settings.

Early work on ViZDoom demonstrated the feasibility of end-to-end learning in first-person environments. (Lample and Chaplot, 2017) showed that deep reinforcement learning agents could learn to navigate, acquire weapons, and engage enemies through trial and error, developing emergent strategies including ambush tactics and hit-and-run maneuvers.

(Dosovitskiy and Koltun, 2017) extended this work to multi-agent settings, studying how agents learned to compete in deathmatch scenarios. The research revealed interesting dynamics including the emergence of territorial behavior and the development of simple communication protocols through in-game actions.

5.4.2 DeepMind Lab and 3D Navigation

DeepMind Lab provided a more sophisticated 3D environment for studying navigation and coordination (Beattie et al., 2016). The platform included procedurally generated mazes, multiple agents, and various task structures enabling systematic study of multi-agent behavior in three-dimensional spaces.

(Jaderberg et al., 2019) used DeepMind Lab to study population-based training in first-person environments, demonstrating that diverse training populations produced more robust and adaptable agents. The research highlighted the importance of opponent diversity for generalization, with agents trained against varied populations performing better against novel opponents.

5.4.3 Counter-Strike: Modern Benchmarks

Recent work has established Counter-Strike: Global Offensive and its successor Counter-Strike 2 as challenging benchmarks for multi-agent reinforcement learning (Pearce et al., 2023). The game's round-based structure, economic system, and emphasis on coordinated team play create unique research opportunities.

The CS:GO benchmark introduced by (Pearce et al., 2023) includes scenarios ranging from simple aim duels to full five-versus-five matches with realistic economics. Initial results revealed the difficulty of learning effective coordination, with current algorithms struggling to match the performance of rule-based agents in complex tactical situations.

(Smith et al., 2024) studied emergent communication in Counter-Strike environments, showing that agents could learn to coordinate through in-game actions such as grenade throws and movement patterns. The research demonstrated that communication need not be explicit; coordinated behavior can emerge from shared understanding of tactical situations.

5.5 Open-World Sandboxes: Minecraft and Open-Ended Learning

Minecraft represents the frontier of multi-agent game environments, offering virtually unlimited possibilities for exploration, creation, and interaction. The game's open-ended nature tests agents' ability to learn diverse skills, set their own goals, and adapt to novel situations.

5.5.1 Project Malmo and the MineRL Competition

Project Malmo provided the first programmable interface to Minecraft, enabling researchers to design custom tasks and multi-agent scenarios (Johnson et al., 2016). The platform supported both single-agent and multi-agent experiments, including cooperative building tasks, competitive survival scenarios, and mixed-motive resource collection games.

The MineRL competition, launched in 2019, challenged participants to develop sample-efficient reinforcement learning algorithms capable of learning complex Minecraft tasks from limited human demonstrations (Guss et al., 2019). The competition dataset included thousands of hours of human gameplay, providing a rich resource for imitation learning and reinforcement learning research.

The competition tasks required agents to obtain diamonds, a multi-step process involving tree chopping, wooden pickaxe crafting, stone mining, iron smelting, and finally diamond mining. This long horizon task with sparse rewards proved

extremely challenging for pure reinforcement learning approaches, driving innovation in hierarchical methods and imitation learning.

5.5.2 Voyager: Open-Ended Exploration with Large Language Models

Voyager introduced a fundamentally new paradigm for open-ended learning in Minecraft by integrating large language models with reinforcement learning (Wang et al., 2023). The system maintained a growing skill library of executable programs, each representing a reusable capability such as "craft wooden pickaxe" or "find animals." A large language model served as the high-level controller, selecting and combining skills to accomplish open-ended goals.

The Voyager architecture consisted of three components: an automatic curriculum that proposed increasingly challenging tasks based on the agent's current capabilities, a skill library that stored and retrieved executable programs, and an iterative prompting mechanism that refined skills through trial and error. The agent continuously explored the Minecraft world, discovering new resources, crafting new items, and expanding its skill library.

Voyager demonstrated capabilities far beyond previous Minecraft agents, including discovering 63 distinct items, crafting diamond tools, and exploring diverse biomes without any human intervention. The system's success highlighted the potential of combining large language models' world knowledge with reinforcement learning's trial-and-error optimization.

5.5.3 Multi-Agent Minecraft: Coordination and Collaboration

Recent work has extended open-ended learning to multi-agent settings in Minecraft. (Fan et al., 2024) introduced a framework for language-guided multi-agent coordination, where agents communicate through natural language to accomplish collaborative building tasks. The system enabled agents to decompose complex construction projects into subtasks, allocate work among team members, and coordinate their efforts through dialogue.

The research revealed interesting dynamics in multi-agent open-ended learning. Agents developed specialized roles based on their initial experiences, with some focusing on resource gathering and others on construction. Communication protocols emerged that balanced expressiveness with efficiency, with agents learning to compress information into concise instructions.

(Chen et al., 2024) studied the emergence of social norms in multi-agent Minecraft societies, showing that populations of agents could develop conventions around resource sharing, territory claims, and division of labor. These norms emerged from individual learning objectives rather than explicit design, demonstrating the potential for bottom-up emergence of social order.

5.6 Card Games: Imperfect Information and Probabilistic Reasoning

Card games present unique challenges through their inherent uncertainty and the need for probabilistic reasoning about hidden information. These environments test agents' ability to maintain beliefs, update probabilities based on observations, and balance exploration with exploitation.

5.6.1 Poker: Pluribus and Equilibrium Finding

Poker has served as the canonical benchmark for imperfect-information games, with research progressing from two-player limit hold'em to six-player no-limit variants. (Brown and Sandholm, 2019) achieved a landmark result with Pluribus, a superhuman artificial intelligence for six-player no-limit Texas hold'em.

Pluribus combined blueprint strategy computation through self-play with real-time search that operated in information space rather than state space. The system maintained a belief distribution over opponent hands and updated this distribution based on betting patterns. During play, Pluribus performed limited-depth search using a novel abstraction technique that reduced the complexity of the game tree while preserving strategic essentials.

The resulting agent demonstrated fundamentally different playing style from human experts, mixing aggressive and passive lines in ways that opponents found difficult to exploit. Pluribus's success validated the feasibility of equilibrium approximation in large-scale imperfect-information games and established techniques that would influence subsequent work.

5.6.2 Hearthstone and Magic: The Gathering

Collectible card games like Hearthstone and Magic: The Gathering introduce additional complexity through large card pools, deck-building mechanics, and random elements. These games test agents' ability to adapt to novel card combinations and reason about probabilistic outcomes.

(Zhang et al., 2022) studied multi-agent learning in Hearthstone Battlegrounds, an auto-battler mode where eight players compete through sequential drafting and combat phases. The research revealed the importance of opponent modeling in environments with hidden information, as successful strategies required anticipating opponents' compositions and adapting accordingly.

(Chen et al., 2023) developed a reinforcement learning framework for Magic: The Gathering that incorporated deck adaptation, enabling agents to modify their strategies based on opponent tendencies. The system demonstrated that meta-learning across matches could improve performance compared to static policies.

5.7 Lessons Learned and Cross-Cutting Insights

The progression of milestones across game genres reveals several cross-cutting insights about multi-agent reinforcement learning.

First, the importance of diverse training opponents emerges consistently across domains. Whether through league training in AlphaStar, population-based training in OpenAI Five, or self-play in AlphaZero, exposure to varied strategies proves essential for robust performance. Agents trained against narrow opponent sets develop exploitable weaknesses and fail to generalize.

Second, the integration of search with learning provides compounding benefits. AlphaGo’s combination of neural networks with Monte Carlo tree search, Pluribus’s information-space search, and Voyager’s iterative prompting all demonstrate that learned representations can guide efficient search, while search can compensate for learned approximations.

Third, representation design critically impacts performance. The hierarchical action spaces in StarCraft II, the attention mechanisms in MOBA agents, and the skill libraries in Voyager all reflect the need for structured representations that capture the compositional nature of complex tasks.

Fourth, the frontier continues to advance toward open-endedness. From fixed board games to procedurally generated worlds, from single tasks to open-ended exploration, the trajectory points toward agents that can continuously adapt and improve without human intervention.

These lessons inform current research directions and suggest promising avenues for future investigation. The next section examines the challenges that remain and the open problems that define the research frontier.

6 Challenges and Open Problems

Despite the remarkable progress documented in the preceding sections, the field of Game-Based Multi-Agent Reinforcement Learning faces numerous fundamental challenges that limit current capabilities and define the research frontier. These challenges span theoretical foundations, algorithmic design, computational scalability, and practical deployment. This section provides a comprehensive examination of the most significant open problems, organized thematically to highlight their interconnections and relative importance.

6.1 Sample Efficiency

The sample efficiency problem represents perhaps the most critical barrier to broader application of multi-agent reinforcement learning. Current algorithms require millions or billions of environment interactions to achieve competent performance, far exceeding the sample efficiency of human learning. This limitation is particularly acute in game environments where each episode may last minutes or hours, and where real-world deployment may be impractical due to time or cost constraints.

6.1.1 Magnitude of the Problem

The sample complexity of state-of-the-art multi-agent reinforcement learning algorithms is staggering. AlphaStar required approximately 200 years of gameplay experience distributed across thousands of parallel agents to achieve grandmaster level (Vinyals et al., 2019). OpenAI Five consumed 180 years of experience per day during training, accumulating tens of millennia of total gameplay (Berner et al., 2019). Even the relatively focused SMAC benchmark typically requires millions of timesteps per scenario, with challenging tasks demanding hundreds of millions of steps for convergence (Samvelyan et al., 2019).

This sample inefficiency stands in stark contrast to human learning. A professional StarCraft player may require thousands of games to reach grandmaster level, representing perhaps two orders of magnitude less experience than AlphaStar. A human learning a new SMAC scenario can achieve competent coordination within dozens of attempts, leveraging prior knowledge about unit capabilities and tactical concepts. This gap suggests fundamental differences in how humans and artificial agents generalize from experience.

6.1.2 Sources of Inefficiency

Several factors contribute to the sample inefficiency of current algorithms. The exploration problem in multi-agent settings is exponentially more complex than in single-agent domains, as agents must jointly discover coordinated strategies. Random exploration in the joint action space becomes infeasible beyond small numbers of agents, necessitating more structured exploration mechanisms.

The non-stationarity problem exacerbates sample inefficiency by rendering past experience potentially obsolete as other agents’ policies change. Experience replay buffers, which provide crucial sample efficiency gains in single-agent settings, must be carefully managed in multi-agent contexts to avoid learning from outdated trajectories (Foerster et al., 2017). Algorithms that discard old experience may waste samples, while those that retain it may learn from irrelevant data.

Credit assignment delays compound these difficulties. In complex game environments, the consequences of individual actions may only become apparent hundreds or thousands of timesteps later. A poorly positioned scout in early-game StarCraft may lead to defeat thirty minutes later, long after the original action has been forgotten by the learning system. Propagating credit across such extended horizons remains extremely challenging.

6.1.3 Promising Directions

Several research directions offer potential pathways toward improved sample efficiency. Model-based reinforcement learning approaches aim to learn environment dynamics and use imagined trajectories for planning, potentially reducing the need for real environment interactions. (Zhang et al., 2021) introduced a multi-agent world model that learned to simulate the dynamics of cooperative games, enabling policy optimization in imagination. The approach achieved comparable performance to

model-free methods with significantly fewer environment interactions.

Imitation learning from human demonstrations provides another avenue for improving sample efficiency. The MineRL competition specifically challenged participants to learn from limited human data, driving innovation in combining imitation with reinforcement learning (Guss et al., 2019). (Baker et al., 2022) showed that video pretraining on large-scale human gameplay data could bootstrap effective policies with minimal subsequent reinforcement learning.

Hierarchical methods address sample efficiency by learning reusable skills that compose into complex behaviors. (Andreas et al., 2017) introduced modular policy architectures that enable transfer between tasks, while (Vezhnevets et al., 2017) developed feudal networks that learn temporal abstractions automatically. These approaches reduce the effective horizon of learning problems and enable credit assignment across longer timescales.

6.1.4 Open Questions

Fundamental questions about sample efficiency remain unresolved. What is the theoretical minimum sample complexity for multi-agent coordination problems? How does this scale with the number of agents, the complexity of interactions, and the horizon length? Can we develop algorithms whose sample complexity grows polynomially rather than exponentially in these parameters?

The relationship between exploration and sample efficiency requires deeper investigation. How should agents explore jointly to efficiently discover coordinated strategies? What role should intrinsic motivation and curiosity play in multi-agent exploration? Can we design exploration strategies that scale gracefully to many agents?

6.2 Generalization

Generalization the ability to perform well in novel situations distinct from the training distribution represents a critical challenge for multi-agent reinforcement learning. Current agents often fail catastrophically when faced with even minor variations in environment conditions, opponent strategies, or task specifications.

6.2.1 Types of Generalization

Generalization in multi-agent game environments takes multiple forms. Zero-shot generalization requires agents to perform well in unseen scenarios without additional training. This includes generalization to new maps, new opponent strategies, new team compositions, and new game modes. Current agents struggle with even modest variations; a StarCraft agent trained on one map may perform poorly on a different map layout, even when all other game mechanics remain identical (Samvelyan et al., 2019).

Few-shot adaptation requires agents to quickly adjust to new situations with limited additional experience. Humans can adapt to new opponents within a single game, adjusting their strategies based on observed tendencies. Current multi-agent reinforcement learning agents typically require extensive retraining to adapt to novel opponents, limiting their practical utility.

Systematic generalization tests whether agents can combine known concepts in novel ways. An agent that has learned to coordinate with zealots and dragoons separately should ideally coordinate with mixed compositions without additional training. Current methods show limited systematic generalization, often failing when familiar elements appear in unfamiliar combinations (Cobbe et al., 2019).

6.2.2 The Generalization Gap

Empirical studies have documented substantial generalization gaps across multiple domains. (Justesen et al., 2018) evaluated reinforcement learning agents on procedurally generated levels, finding that performance degraded significantly on unseen level configurations. (Zhang et al., 2018) studied generalization in first-person shooter environments, showing that agents trained on fixed maps failed to transfer to new map layouts even when visual appearances were similar.

The SMAC benchmark includes a generalization protocol where agents trained on specific unit compositions are tested on variations with different numbers of units or altered unit attributes (Samvelyan et al., 2019). Results consistently show performance degradation, with some algorithms losing over 50% of their win rate on held-out configurations. This generalization gap persists even when training includes extensive randomization, suggesting fundamental limitations in current learning approaches.

6.2.3 Sources of Poor Generalization

Several factors contribute to poor generalization in multi-agent reinforcement learning. Overfitting to training opponents represents a primary cause; agents learn to exploit specific weaknesses rather than develop robust strategies. In league training, agents may converge to strategies that beat the current population but fail against out-of-distribution opponents (Vinyals et al., 2019).

Observation overfitting occurs when agents rely on spurious correlations in the observation space. An agent might learn to associate a particular visual feature with enemy presence, only to fail when that feature appears in different contexts. This problem is exacerbated by the high-dimensional observations typical of game environments.

Policy specialization emerges when agents learn complementary but brittle coordination patterns. A team of agents may develop a coordinated strategy that works well in training but fails when any agent deviates slightly from expected behavior. This brittleness limits robustness to novel situations where perfect coordination may be impossible.

6.2.4 Approaches to Improving Generalization

Several techniques have shown promise for improving generalization. Domain randomization exposes agents to diverse training conditions, encouraging policies that work across variations. (Tobin et al., 2017) introduced domain randomization for robotics, and subsequent work has applied similar ideas to game environments. (Cobbe et al., 2019) showed that training on procedurally generated levels substantially improved generalization in reinforcement learning benchmarks.

Regularization techniques adapted from supervised learning can reduce overfitting in multi-agent settings. (Farebrother et al., 2018) applied L2 regularization and dropout to deep Q-networks, finding modest improvements in generalization. (Igl et al., 2019) introduced information bottleneck regularization that encourages policies to ignore task-irrelevant features.

Population-based training naturally promotes generalization by exposing agents to diverse opponents. The league training framework in AlphaStar generated robust strategies through exposure to varied play styles (Vinyals et al., 2019). (Czarnecki et al., 2020) formalized this insight, showing that diverse training populations produce policies that generalize better to novel opponents.

Meta-learning offers a framework for learning to adapt quickly to new situations. (Al-Shedivat et al., 2018) introduced meta-learning for multi-agent reinforcement learning, enabling agents to adapt to new opponents within episodes. (Kim et al., 2021) extended this approach to team coordination, showing that meta-learned policies could adapt to novel teammates.

6.2.5 Open Questions

The generalization problem raises fundamental questions about the nature of learning in multi-agent systems. What constitutes a proper test of generalization? How should we measure progress beyond performance on held-out instances from the same distribution? Can we develop benchmarks that systematically test different forms of generalization?

The relationship between representation learning and generalization requires deeper investigation. What representations support systematic generalization? How can we encourage agents to learn causal structures rather than statistical correlations? Can unsupervised pretraining on diverse data provide foundations for generalization?

6.3 Scalability

Scalability challenges permeate multi-agent reinforcement learning, affecting algorithmic complexity, computational requirements, and the ability to handle large numbers of agents. These challenges become increasingly acute as the field progresses toward environments with hundreds or thousands of interacting agents.

6.3.1 Curse of Dimensionality

The joint action space grows exponentially with the number of agents, creating fundamental combinatorial challenges. For n agents each with $|\mathcal{A}|$ actions, the number of possible joint actions is $|\mathcal{A}|^n$, rendering exhaustive consideration infeasible beyond small n . This curse of dimensionality affects planning, exploration, and credit assignment.

Value-based methods address this challenge through factorization, but the representational capacity of factorized value functions remains incompletely understood. (Mahajan et al., 2021) analyzed the representational power of monotonic value factorization, showing that while any function can be represented in principle, practical constraints limit expressiveness. The trade-off between representational capacity and tractable learning remains an active research area.

Policy gradient methods avoid explicit representation of the joint action space but face increasing gradient variance as the number of agents grows. (Papoudakis et al., 2020) empirically characterized this variance growth, showing that baseline techniques become less effective with many agents. The theoretical scaling of gradient variance with agent count remains poorly understood.

6.3.2 Computational Scaling

The computational demands of multi-agent reinforcement learning grow rapidly with environment complexity and agent count. Training AlphaStar required thousands of TPUs for weeks, representing computational resources beyond the reach of most research groups (Vinyals et al., 2019). OpenAI Five utilized tens of thousands of CPUs and hundreds of GPUs (Bernier et al., 2019). This computational intensity creates barriers to entry and limits the pace of research.

Communication overhead becomes significant in large-scale multi-agent systems. Algorithms that require sharing gradients or value estimates between agents face bandwidth limitations as the number of agents grows. Decentralized training paradigms reduce communication requirements but may sacrifice coordination quality.

The memory requirements for storing experience replay buffers scale with the number of agents and the trajectory length. For environments with hundreds of agents and long episodes, experience storage becomes prohibitively expensive. This limitation has driven interest in on-policy methods that avoid replay buffers, though these may be less sample efficient.

6.3.3 Many-Agent Challenges

Environments with hundreds or thousands of agents introduce qualitatively new challenges beyond those present in small-scale systems. Neural MMO supports up to 1024 agents in persistent virtual worlds, enabling study of large-scale multi-agent dynamics (Suarez et al., 2019). The resulting behaviors exhibit emergent properties including resource competition, territoriality, and population-level specialization.

Coordination in large agent populations requires different mechanisms than in small teams. Pairwise coordination mechanisms become intractable as the number of potential interactions grows quadratically. Agents must learn to identify relevant

interaction partners and ignore irrelevant others. Attention mechanisms provide one approach, but their computational cost scales with the number of agents considered.

Emergent division of labor in large populations creates new challenges for credit assignment. When hundreds of agents contribute to collective outcomes, determining individual contributions becomes extremely difficult. Population-level metrics may obscure important variations in individual behavior.

6.3.4 Promising Directions

Several research directions address scalability challenges. Mean-field reinforcement learning approximates many-agent interactions through aggregate statistics, reducing the effective dimensionality of the problem (Yang et al., 2018). The approach assumes that agents are exchangeable and that the impact of any single agent on others can be approximated by the population distribution. While restrictive, this assumption enables scaling to very large populations.

Graph neural networks provide natural representations for multi-agent systems, enabling policies that generalize across different interaction structures (Battaglia et al., 2018). Agents can be represented as nodes in a graph, with edges representing potential interactions. Message passing through the graph enables coordination while maintaining locality.

Decentralized training paradigms reduce communication and coordination overhead. (Zhang et al., 2018) introduced distributed deep reinforcement learning frameworks that scale to thousands of agents through asynchronous training. (Espeholt et al., 2018) developed IMPALA, a scalable architecture for distributed reinforcement learning that has been applied to multi-agent domains.

6.3.5 Open Questions

Fundamental questions about scalability remain unresolved. What is the maximum number of agents that can be effectively coordinated with current methods? How does required coordination quality scale with agent count? Can we develop algorithms whose complexity grows linearly rather than exponentially in agent number?

The relationship between agent specialization and scalability requires investigation. As populations grow, agents may naturally specialize, reducing the need for global coordination. Can we design algorithms that encourage beneficial specialization while maintaining adaptability?

6.4 Credit Assignment

The multi-agent credit assignment problem determining how to attribute global outcomes to individual agent actions remains fundamentally challenging despite significant progress. This problem becomes particularly acute in environments with delayed rewards, complex causal chains, and heterogeneous agent capabilities.

6.4.1 Temporal Credit Assignment

Temporal credit assignment the problem of determining which past actions caused current rewards becomes more difficult in multi-agent settings due to the interdependence of agent actions. A successful team fight may result from a sequence of actions by multiple agents, with complex interactions between ability timing, positioning, and target selection. Disentangling these causal relationships requires reasoning about counterfactuals.

Value decomposition methods address spatial credit assignment across agents but provide limited mechanisms for temporal credit assignment. The Bellman update propagates value information backward through time, but this process may be slow when rewards are delayed by many steps. Eligibility traces provide one mechanism for accelerating temporal credit assignment, but their effectiveness in multi-agent settings remains underexplored (Schroeder de Witt et al., 2020).

Hierarchical methods offer a complementary approach by learning skills at multiple timescales. A high-level policy selects skills that persist for many timesteps, enabling credit assignment at the skill level rather than the primitive action level. Tang2018 demonstrated improved credit assignment in hierarchical multi-agent systems, but the automatic discovery of useful skills remains challenging.

6.4.2 Spatial Credit Assignment

Spatial credit assignment across agents has received substantial attention through value decomposition methods. VDN, QMIX, and their variants provide mechanisms for factorizing global value functions into agent-specific components (Sunehag et al., 2017; Rashid et al., 2018). However, these methods assume particular factorization structures that may not match the true causal structure of the environment.

Counterfactual methods offer a more direct approach to credit assignment by comparing actual outcomes to counterfactual baselines. COMA computes advantages by marginalizing out individual agent actions, providing each agent with a measure of its unique contribution (Foerster et al., 2018). The computational cost of counterfactual evaluation scales with the action space size, limiting applicability to domains with large action spaces.

Difference rewards provide another mechanism for credit assignment by measuring the change in global reward when an agent's action is replaced with a default action (Wolpert and Tumer, 2001). This approach requires the ability to simulate counterfactual outcomes, which may be available in some game environments through deterministic simulation or learned world models.

6.4.3 Structural Credit Assignment

Beyond temporal and spatial dimensions, credit assignment involves structural questions about how agent interactions create value. Some agent actions may enable others' actions without directly contributing to reward. A scout that reveals enemy positions enables subsequent attacks, but the credit for successful attacks may be incorrectly assigned to the attackers alone.

Causal inference methods offer potential solutions to structural credit assignment by learning causal models of multi-agent interactions. (Jaques et al., 2019) introduced social influence metrics that measure an agent's impact on others' behavior, providing a mechanism for attributing credit through social effects. (Wang et al., 2020) developed counterfactual policy evaluation methods that estimate the causal effect of agent actions on team outcomes.

6.4.4 Open Questions

The credit assignment problem raises fundamental questions about the nature of multi-agent learning. What is the appropriate credit assignment target for cooperative learning? Should agents receive credit proportional to their causal contribution, or should credit be distributed according to fairness principles? How do different credit assignment mechanisms affect the emergence of cooperation and specialization?

The relationship between credit assignment and exploration requires investigation. Effective credit assignment enables agents to identify which exploration attempts produced valuable outcomes, accelerating learning. Can we design credit assignment mechanisms that explicitly guide exploration toward high-impact actions?

6.5 Non-Stationarity and Convergence

The non-stationarity inherent in multi-agent learning creates fundamental challenges for theoretical analysis and practical algorithm design. As all agents learn simultaneously, the environment from any agent's perspective changes continuously, violating the assumptions underlying much of reinforcement learning theory.

6.5.1 The Non-Stationarity Problem

In single-agent reinforcement learning, the environment dynamics are stationary, meaning that the transition probabilities and reward functions remain constant throughout learning. This stationarity enables theoretical guarantees including convergence to optimal policies under appropriate conditions. In multi-agent settings, the effective dynamics from any agent's perspective depend on the policies of other agents, which are themselves changing through learning.

This non-stationarity manifests in several ways. An agent's experience replay buffer contains trajectories generated under different opponent policies, potentially mixing incompatible data. Policy gradient estimates may have high variance due to changing opponent behavior. Value functions may become inconsistent as the underlying policy distribution shifts.

(Hernandez-Leal et al., 2019) provided a comprehensive analysis of non-stationarity in multi-agent reinforcement learning, documenting its effects on various algorithm classes. The severity of non-stationarity depends on the learning dynamics; rapid policy updates and heterogeneous learning rates exacerbate the problem.

6.5.2 Convergence Guarantees

Theoretical convergence results for multi-agent reinforcement learning are substantially weaker than for single-agent settings. In general-sum stochastic games, no algorithm guarantees convergence to Nash equilibrium without restrictive assumptions. Even in two-player zero-sum games, convergence may require specific update schedules or equilibrium selection mechanisms.

(Littman, 1994) established convergence for Q-learning in zero-sum stochastic games under the assumption that agents maintain fixed policies during learning. (Bowling and Veloso, 2002) introduced the concept of rational and convergent learning, showing that WoLF policy gradient methods could converge in self-play. However, these results apply to specific algorithm classes and game structures.

Recent work has made progress on convergence guarantees for particular algorithm families. (Pérolat et al., 2015) established asymptotic convergence for gradient-based methods in potential games. (Leonardos and Piliouras, 2021) analyzed the dynamics of policy gradient in zero-sum games, showing connections to Hamiltonian dynamics. Despite these advances, general convergence guarantees remain elusive.

6.5.3 Practical Approaches

Several practical approaches mitigate non-stationarity without providing theoretical guarantees. Centralized training with decentralized execution reduces non-stationarity during learning by providing stable value functions that condition on all agents' actions. The centralized critic in MADDPG and related algorithms remains stationary even as policies change, providing stable targets for policy improvement (Lowe et al., 2017).

Experience replay management techniques can reduce the impact of non-stationarity. (Foerster et al., 2017) introduced importance sampling corrections for multi-agent experience replay, adjusting for the probability that trajectories were generated under different policies. (Palmer et al., 2018) developed lenient learning approaches that forgive outdated experience based on its age.

Population-based methods naturally handle non-stationarity by maintaining diverse policies and adapting to the current population distribution. League training in AlphaStar continuously adjusted to the evolving policy distribution, with main agents facing a non-stationary but manageable opponent set (Vinyals et al., 2019).

6.5.4 Open Questions

The non-stationarity problem raises fundamental questions about the nature of multi-agent learning. Is stationarity a necessary condition for reliable learning, or can we develop algorithms that thrive in non-stationary environments? What are the minimal assumptions required for convergence in multi-agent settings?

The relationship between non-stationarity and exploration requires investigation. Non-stationarity may actually benefit exploration by automatically generating diverse experience as opponent policies change. Can we design algorithms that harness beneficial non-stationarity while mitigating its harmful effects?

6.6 Opponent Modeling and Theory of Mind

Opponent modeling the ability to infer and predict other agents' strategies becomes increasingly important in competitive and mixed-motive environments. However, effective opponent modeling raises profound challenges including recursive reasoning, computational tractability, and the integration of modeling with decision-making.

6.6.1 The Recursion Problem

Opponent modeling naturally leads to recursive reasoning: agent A models agent B, but agent B may also model agent A, and both may model the other's modeling. This recursion can continue indefinitely, creating infinite regress. In game theory, this recursion is resolved through equilibrium concepts where beliefs are consistent with actual play, but learning agents must approximate this consistency online.

Level-k reasoning provides a bounded approximation to recursive modeling. Level-0 agents follow fixed heuristics; level-1 agents best-respond to level-0; level-2 agents best-respond to level-1, and so on. (Camerer, 2004) showed that human reasoning rarely exceeds level-3 in practice, suggesting bounded recursion may be sufficient. Multi-agent reinforcement learning agents have been trained to implement level-k reasoning, achieving improved performance in competitive games (Zhang et al., 2020).

Deep theory of mind approaches learn recursive models through neural architectures that simulate others' reasoning. (Rabinowitz et al., 2018) introduced a machine theory of mind framework that learns to predict others' behavior by simulating their mental states. (Foerster et al., 2019) developed Bayesian opponent modeling that maintains distributions over opponent types and updates beliefs based on observations.

6.6.2 Computational Challenges

Opponent modeling adds substantial computational overhead to multi-agent learning. Maintaining explicit models of all other agents requires memory that scales linearly with agent count, and updating these models requires additional computation. In large-scale environments with many agents, comprehensive opponent modeling becomes infeasible.

Selective attention mechanisms can reduce computational burden by focusing modeling resources on relevant opponents. (Das et al., 2019) introduced attention-based communication that implicitly models message relevance, while (Rabinowitz et al., 2018) learned to predict which agents' behavior was most informative for task performance.

The integration of opponent modeling with decision-making presents additional challenges. Agents must decide how much to rely on opponent models versus learned policies, balancing exploration of opponent behavior with exploitation of known responses. Model uncertainty must be incorporated into decision-making, with agents behaving more cautiously when opponent models are uncertain.

6.6.3 Open Questions

Opponent modeling raises fundamental questions about the nature of strategic reasoning. What level of recursive modeling is necessary for optimal play in various game classes? How does the required modeling depth scale with game complexity and agent count? Can agents learn to dynamically adjust their modeling depth based on task demands?

The relationship between opponent modeling and robustness requires investigation. Accurate opponent models may enable exploitation of specific opponents but could lead to overfitting and brittleness. How can agents balance modeling accuracy with robustness to novel opponents?

6.7 Explainability and Interpretability

As multi-agent reinforcement learning systems become more capable and are deployed in consequential applications, the need for explainability and interpretability grows. Understanding why agents make particular decisions, how they coordinate, and what strategies they have learned is essential for trust, debugging, and safety assurance.

6.7.1 The Black Box Problem

Deep neural network policies are notoriously difficult to interpret, and this opacity is compounded in multi-agent settings where behavior emerges from interactions among multiple black boxes. A team of agents may achieve high performance through coordination patterns that are opaque to human observers, making it difficult to verify that behavior aligns with intended objectives.

The SMAC benchmark illustrates this opacity. While algorithms can achieve high win rates, understanding exactly how they coordinate whether they use focus fire, ability combos, or positioning strategies requires extensive post-hoc analysis

(Samvelyan et al., 2019). Different algorithms may achieve similar performance through qualitatively different coordination patterns, but these differences are not captured by aggregate metrics.

6.7.2 Explainability Methods

Several approaches have been developed to explain multi-agent reinforcement learning behavior. Saliency maps highlight input features that influence agent decisions, revealing which observations agents attend to (Greydanus et al., 2018). In StarCraft, saliency analysis shows which enemy units influence agent targeting decisions, providing insight into coordination patterns.

Behavioral analysis through controlled experiments can reveal coordination mechanisms. By systematically varying environmental conditions and observing resulting behavior, researchers can infer strategic principles underlying agent policies. (Rabiner et al., 2021) used this approach to analyze cooperative behavior in Hanabi, identifying the emergence of implicit communication protocols.

Natural language explanations generated by large language models offer a promising direction for interpretability. (Park et al., 2023) showed that language models could generate plausible explanations of agent behavior when provided with observation-action trajectories. (Zhang et al., 2024) extended this approach to multi-agent settings, generating dialogue that explained team coordination strategies.

6.7.3 Open Questions

Explainability raises fundamental questions about the nature of understanding in multi-agent systems. What constitutes an adequate explanation of multi-agent behavior? How can we verify that explanations faithfully represent agent reasoning rather than plausible post-hoc justifications? Can we design agents that are inherently more interpretable without sacrificing performance?

The relationship between explainability and trust requires investigation. Do explanations actually increase human trust in multi-agent systems, and if so, what forms of explanation are most effective? How can explanations support debugging and improvement of multi-agent reinforcement learning systems?

6.8 Safety and Robustness

As multi-agent reinforcement learning systems approach deployment in real-world applications, safety and robustness considerations become paramount. Game environments provide controlled settings for studying these issues, which have direct implications for autonomous driving, financial trading, and multi-robot systems.

6.8.1 Distributional Shift and Robustness

Multi-agent reinforcement learning agents trained in simulation may encounter distributional shift when deployed in real environments or even in different game scenarios. This shift can lead to catastrophic failures if agents encounter situations outside their training distribution.

Adversarial attacks on multi-agent policies pose particular concerns. Small perturbations to agent observations can induce dramatically different actions, potentially causing coordination failures or unsafe behavior (Gleave et al., 2020). In competitive games, adversaries might exploit these vulnerabilities to manipulate agent behavior.

Robustness to teammate failures becomes critical in cooperative settings. Agents trained with the expectation that all teammates will act optimally may fail catastrophically when teammates malfunction or behave suboptimally. Training with variable teammate capabilities can improve robustness but may reduce performance with optimal teammates (Carroll et al., 2019).

6.8.2 Safe Exploration

Exploration in multi-agent systems carries risks that are absent in single-agent settings. An agent exploring novel strategies may inadvertently cause harm to teammates or opponents, or may trigger cascading failures through unexpected interactions. In game environments, these risks are contained, but the methods developed for safe exploration may transfer to real-world applications.

Constrained reinforcement learning frameworks provide mechanisms for incorporating safety constraints into multi-agent learning. (Chow et al., 2018) developed safe policy optimization algorithms that maintain constraint satisfaction throughout learning. (García and Fernández, 2015) surveyed safe exploration methods, many of which extend naturally to multi-agent settings.

6.8.3 Alignment and Value Learning

The alignment problem ensuring that agent objectives align with human values becomes more complex in multi-agent settings. Even if individual agents are aligned, their interactions may produce collectively undesirable outcomes. The tragedy of the commons exemplifies this challenge: individually rational agents may collectively destroy shared resources.

Inverse reinforcement learning from human demonstrations can infer reward functions that capture human preferences (Ng and Russell, 2000). In multi-agent settings, this requires inferring both individual and collective preferences from observed behavior. (Hadfield-Menell et al., 2016) introduced cooperative inverse reinforcement learning, where agents actively query humans to resolve reward uncertainty.

6.8.4 Open Questions

Safety in multi-agent systems raises fundamental questions about the nature of robust coordination. How can we guarantee that multi-agent policies remain safe under all possible opponent behaviors? What are the minimal assumptions required for safety guarantees in multi-agent settings? How can we design reward functions that encode safety constraints without inducing reward hacking?

The relationship between safety and exploration requires investigation. How can agents explore safely in multi-agent environments where the consequences of exploration depend on others' responses? Can we develop theoretical frameworks for safe multi-agent exploration analogous to those in single-agent settings?

6.9 Evaluation and Benchmarking

Rigorous evaluation of multi-agent reinforcement learning algorithms presents unique challenges that complicate progress measurement and comparison across approaches. These challenges span statistical significance, reproducibility, and the design of meaningful benchmarks.

6.9.1 Statistical Challenges

Multi-agent reinforcement learning evaluation suffers from high variance due to the stochasticity of both environment dynamics and opponent policies. A single run may produce widely varying results, requiring many seeds for statistically meaningful comparisons. (Colas et al., 2018) documented the scale of this variance in common benchmarks, showing that hundreds of seeds may be necessary to detect moderate improvements.

The non-stationarity of evaluation compounds these difficulties. As opponents improve through learning, the distribution of test cases changes over time, complicating comparisons between algorithms trained for different durations. Fixed opponent baselines provide stable evaluation but may not reflect performance against learning opponents.

6.9.2 Benchmark Design

Effective benchmarks must balance multiple desiderata: they should be challenging enough to drive progress, accessible enough to enable broad participation, and structured to provide meaningful comparisons. The SMAC benchmark has succeeded in many respects, becoming a standard testbed for cooperative multi-agent reinforcement learning (Samvelyan et al., 2019). However, concerns have been raised about overfitting to specific scenarios and the ecological validity of results.

The Melting Pot benchmark addresses these concerns by providing diverse social dilemmas and measuring generalization across scenarios (Agapiou et al., 2022). The benchmark includes multiple game structures, enabling evaluation of social intelligence beyond narrow task performance. Early results reveal substantial variation in algorithm performance across scenarios, suggesting that current methods lack general social reasoning capabilities.

6.9.3 Reproducibility

Reproducibility challenges in multi-agent reinforcement learning are exacerbated by the complexity of implementations and the sensitivity of results to hyperparameters. (Henderson et al., 2018) documented widespread reproducibility failures in reinforcement learning research, and multi-agent settings compound these difficulties.

Open-source implementations and standardized evaluation protocols can improve reproducibility. Frameworks such as PyMARL provide unified implementations of multiple algorithms, enabling fair comparisons (Samvelyan et al., 2019). However, even with shared codebases, subtle implementation details can substantially affect results.

6.9.4 Open Questions

Evaluation raises fundamental questions about progress measurement in multi-agent reinforcement learning. What metrics appropriately capture multi-agent capabilities? How can we measure generalization, robustness, and sample efficiency alongside raw performance? What role should human evaluation play in assessing progress?

The relationship between benchmark performance and real-world capability requires investigation. Do improvements on game benchmarks translate to improved performance in real-world multi-agent systems? How can we design benchmarks that better predict deployment performance?

The challenges outlined in this section define the research frontier of Game-Based Multi-Agent Reinforcement Learning. Progress on any of these fronts would represent substantial advancement, and the most significant breakthroughs will likely address multiple challenges simultaneously. The following section examines promising future directions that may lead to such breakthroughs.

7 Conclusions and Future Directions

The intersection of digital games and multi-agent reinforcement learning has emerged as one of the most vibrant and consequential research areas in contemporary artificial intelligence. This comprehensive review has traced the evolution of Game-MARL from its theoretical foundations in stochastic games to the current frontier of open-ended learning in complex virtual worlds. We have examined the core algorithms that drive progress in the field, the landmark applications that have demonstrated the capabilities of multi-agent systems, and the fundamental challenges that define the research frontier. In this concluding section, we synthesize the key insights emerging from this review and chart promising directions for future investigation.

7.1 Summary of Key Insights

The trajectory of Game-MARL research reveals several overarching themes that characterize the field's development and inform its future direction.

7.1.1 The Co-Evolution of Environments and Algorithms

Throughout the history of Game-MARL, advances in algorithms and environments have progressed in tandem, each driving the other forward. Early work on simple grid-world scenarios and classic board games established the foundational frameworks of stochastic games and multi-agent Q-learning (Littman, 1994). The introduction of the Arcade Learning Environment and Atari games catalyzed the deep reinforcement learning revolution, demonstrating that neural networks could learn directly from high-dimensional sensory inputs (Mnih et al., 2013, 2015). Subsequent environments of increasing complexity including StarCraft II, Dota 2, and Minecraft have pushed algorithms to develop increasingly sophisticated coordination mechanisms, hierarchical representations, and population-based training paradigms (Vinyals et al., 2019; Berner et al., 2019; Guss et al., 2019).

This co-evolutionary dynamic shows no signs of slowing. Each new generation of environments reveals limitations in current algorithms, driving innovation that in turn enables progress toward more ambitious domains. The recent emergence of open-ended learning environments like Neural MMO and Melting Pot signals the next frontier, challenging researchers to develop algorithms capable of open-ended adaptation and social reasoning (Suarez et al., 2019; Agapiou et al., 2022).

7.1.2 The Primacy of Coordination Mechanisms

A central insight emerging from decades of research is that effective multi-agent learning requires explicit or implicit mechanisms for coordination. Value decomposition methods address this need through factorized value functions that enable agents to evaluate joint actions while maintaining decentralized execution (Sunehag et al., 2017; Rashid et al., 2018). Centralized training with decentralized execution frameworks provide stable learning signals through critics that condition on global information (Lowe et al., 2017; Foerster et al., 2018; Yu et al., 2022). Communication-based methods enable explicit information sharing, with architectures ranging from simple message broadcasting to attention-based selective communication (Sukhbaatar et al., 2016; Das et al., 2019; Singh et al., 2019).

The diversity of coordination mechanisms reflects the variety of coordination requirements across different game genres. Cooperative micromanagement in StarCraft II benefits from value decomposition, while mixed-motive environments like Diplomacy require more sophisticated reasoning about intentions and commitments (Bakhtin et al., 2022). The field's growing appreciation for this diversity has led to more nuanced evaluations and algorithm selection guided by problem characteristics rather than one-size-fits-all solutions.

7.1.3 The Power of Population-Based Training

The development of population-based training methods represents a paradigm shift in how multi-agent systems are trained. Rather than learning against fixed opponents or simple self-play, modern systems maintain diverse populations of agents that provide rich and varied training experiences (Vinyals et al., 2019; Berner et al., 2019). League training frameworks with specialized agent roles main agents, league explorers, and past checkpoints generate automatic curricula that continuously challenge agents to improve.

Theoretical analyses have begun to illuminate why population-based methods succeed. Diverse training populations prevent convergence to narrow strategies that exploit specific weaknesses, promoting robustness and generalization (Czarnecki et al., 2020). The competitive dynamics within populations create natural pressure toward more sophisticated behaviors, with agents needing to continuously adapt to survive. These insights suggest that the social dimension of multi-agent learning the interaction within populations of co-evolving agents may be as important as the interaction with the environment itself.

7.1.4 The Centrality of Representation

Across all domains and algorithm families, representation design emerges as a critical determinant of success. The hierarchical action spaces in StarCraft II, the attention mechanisms in MOBA agents, the transformer architectures in modern value decomposition methods, and the skill libraries in open-ended learning systems all reflect the need for structured representations that capture the compositional nature of complex tasks (Vinyals et al., 2019; Wu et al., 2021; Wen et al., 2022; Wang et al., 2023).

The choice of representation affects not only performance but also sample efficiency, generalization, and interpretability. Representations that align with the underlying structure of the task enable more efficient learning and better transfer to novel situations. The growing integration of large language models into multi-agent systems represents a new frontier in representation learning, leveraging vast pre-trained knowledge to bootstrap understanding of game mechanics and strategic concepts (Wang et al., 2023; Fan et al., 2024).

7.2 Future Directions

Building on the insights emerging from this review, we identify several promising directions for future research that address current limitations and push toward more capable multi-agent systems.

7.2.1 Foundation Models for Multi-Agent Systems

The integration of large language models and other foundation models into multi-agent reinforcement learning represents perhaps the most exciting frontier in the field. These models bring extensive world knowledge, few-shot learning capabilities, and natural language understanding that could fundamentally transform how agents learn and coordinate.

Several promising directions warrant exploration. Language models could serve as high-level planners that decompose complex tasks into sequences of simpler skills, as demonstrated in Voyager for single-agent Minecraft (Wang et al., 2023). Extending this approach to multi-agent settings would enable teams of agents to coordinate through natural language, negotiating roles, sharing information, and adapting plans based on changing circumstances (Fan et al., 2024).

Foundation models could also provide rich priors for multi-agent learning, reducing the sample complexity of acquiring basic competencies. Pre-trained representations of game concepts, strategic principles, and common-sense physics could bootstrap learning in new environments, enabling agents to achieve competent performance with orders of magnitude less experience than current methods require.

The challenge of grounding language models in interactive experience remains significant. Language models trained on static text lack the causal understanding that comes from acting in environments and observing consequences. Developing methods that enable language models to learn from interaction while retaining their broad knowledge base is a key open problem.

7.2.2 Open-Ended Learning and Auto-Curricula

The progress toward open-ended learning systems that continuously generate their own training tasks represents a fundamental shift in how we think about multi-agent intelligence. Rather than training on fixed distributions of scenarios, open-ended systems explore open worlds, discover new challenges, and develop increasingly sophisticated capabilities over unbounded time horizons.

The Voyager system demonstrated the potential of open-ended learning in single-agent settings, with the agent continuously exploring Minecraft and expanding its skill library without human intervention (Wang et al., 2023). Extending this paradigm to multi-agent settings raises fascinating questions about the emergence of specialization, trade, and social organization. Populations of agents exploring open worlds might naturally develop division of labor, resource exchange, and collective decision-making mechanisms.

Auto-curricula that generate increasingly challenging tasks based on current agent capabilities could accelerate progress toward general intelligence. Rather than relying on human-designed benchmarks, future systems might generate their own evaluation tasks, ensuring continuous pressure toward improvement. The challenge lies in designing curricula that promote genuine skill acquisition rather than narrow overfitting to generated tasks.

7.2.3 Human-AI Coordination and Teaming

As multi-agent systems approach deployment in real-world applications, the ability to coordinate effectively with humans becomes increasingly important. Human-AI coordination raises unique challenges beyond those present in pure artificial systems, including the need to interpret human intentions, adapt to human variability, and communicate in ways that humans can understand.

The Hanabi challenge has provided a controlled setting for studying human-AI coordination, revealing the importance of theory of mind and common-sense reasoning (Bard et al., 2020). Diplomacy takes this further, requiring natural language negotiation and alliance formation (Bakhtin et al., 2022). Future work should extend these investigations to more complex domains, developing methods for agents to infer human preferences from limited interactions and to adapt their behavior accordingly.

Reciprocity and trust mechanisms will be essential for effective human-AI teams. Agents must be able to signal their intentions reliably, to honor commitments, and to respond appropriately to human trust and distrust. Learning these social competencies in controlled game environments could provide foundations for deployment in high-stakes real-world settings.

7.2.4 Theoretical Foundations

Despite empirical progress, the theoretical foundations of multi-agent reinforcement learning remain underdeveloped. Rigorous understanding of convergence conditions, sample complexity, and representational requirements would accelerate progress and guide algorithm design.

The relationship between game structure and algorithmic suitability deserves deeper investigation. Which classes of games admit efficient learning with value decomposition? When are policy gradient methods preferable to value-based approaches? How does the required coordination mechanism relate to the payoff structure and information availability? Answering these questions would enable more principled algorithm selection and design.

The sample complexity of multi-agent learning remains poorly understood. Theoretical bounds that characterize how required experience scales with number of agents, horizon length, and task complexity would provide guidance for algorithm development and benchmark design. Understanding the fundamental limits of multi-agent learning could also inform decisions about when to pursue learning-based approaches versus alternative methods.

7.2.5 Beyond Games: Real-World Transfer

The ultimate motivation for Game-MARL research lies in its potential to enable real-world multi-agent systems. Autonomous vehicles coordinating in traffic, robot teams performing search and rescue, and economic agents interacting in markets could

all benefit from advances in multi-agent learning.

However, the gap between game environments and real-world deployment remains substantial. Games provide clean observation spaces, well-defined reward functions, and safe reset mechanisms that are absent in physical systems. Bridging this gap requires progress on multiple fronts: robustness to observation noise, safe exploration in costly environments, and transfer from simulation to reality.

Game environments can serve as proving grounds for techniques that will eventually enable real-world deployment. Work on domain randomization, robust policy learning, and sim-to-real transfer in game contexts could translate directly to robotics and other physical domains (Tobin et al., 2017). The challenge lies in identifying which aspects of game-based research generalize and which are artifacts of particular game structures.

7.2.6 Social and Ethical Implications

As multi-agent systems become more capable, attention to their social and ethical implications becomes increasingly urgent. Game environments provide controlled settings for studying these issues before deployment in high-stakes contexts.

The emergence of cooperation and conflict in multi-agent systems raises questions about alignment with human values. Agents trained purely to maximize game scores may develop strategies that are effective but undesirable from a human perspective perhaps overly aggressive, deceptive, or exploitative. Understanding how to shape agent behavior through reward design and training paradigms is essential for beneficial deployment.

The potential for multi-agent systems to be used in manipulative or adversarial applications also warrants consideration. Game environments could serve as testbeds for studying detection and defense against adversarial multi-agent systems, developing techniques that would translate to real-world security applications.

Finally, the economic and social impacts of capable multi-agent systems deserve careful study. Game environments that simulate economic interactions, resource competition, and social dilemmas can provide insights into how autonomous agents might affect human welfare in real-world contexts (Leibo et al., 2017; Hughes et al., 2018; McKee et al., 2020).

7.3 Concluding Remarks

The journey of Game-Based Multi-Agent Reinforcement Learning from simple grid-worlds to grandmaster-level StarCraft play and open-ended Minecraft exploration spans barely a decade. This rapid progress testifies to the power of games as platforms for artificial intelligence research and to the creativity and persistence of the research community. Yet the challenges that remain are as exciting as the achievements that have been realized.

The path forward requires continued integration of ideas from diverse fields: game theory for understanding strategic interaction, deep learning for rich representation, cognitive science for insights into coordination and communication, and robotics for grounding in physical reality. The game environments that have served as testbeds for progress will continue to evolve, presenting new challenges that push the boundaries of what multi-agent systems can achieve.

For researchers entering the field, the opportunities are abundant. Foundational questions about credit assignment, generalization, and non-stationarity remain unresolved. New frontiers in open-ended learning, human-AI coordination, and foundation-model integration are just beginning to open. The tools and techniques developed in Game-MARL research will increasingly find application beyond games, contributing to the broader goal of artificial intelligence that can understand, interact with, and improve the world.

Digital games have proven to be not merely entertainment but also a mirror for intelligence itself. In the constrained but rich worlds of games, we can observe the emergence of coordination, competition, communication, and cooperation. As our agents become more sophisticated, the games they play and the strategies they discover will continue to illuminate the nature of multi-agent intelligence and its potential to transform our world.

References

- John Agapiou, Alexander S Vezhnevets, Edgar Duenez-Guzman, Jayd Matyas, Yiran Mao, Peter Sunehag, Oliver Köpf, Janosch Kramar, Jonathan Yang, Jane Wang, et al. Melting pot 2.0: A platform for studying multi-agent reinforcement learning and social intelligence. *arXiv preprint arXiv:2211.13746*, 2022.
- Maruan Al-Shedivat, Trapit Bansal, Yuri Burda, Ilya Sutskever, Igor Mordatch, and Pieter Abbeel. Continuous adaptation via meta-learning in nonstationary and competitive environments. *International Conference on Learning Representations*, 2018.
- Stefano V Albrecht and Peter Stone. Autonomous agents modelling other agents: A comprehensive survey and open problems. *Artificial Intelligence*, 258:66–95, 2018.
- Jacob Andreas, Dan Klein, and Sergey Levine. Modular multitask reinforcement learning with policy sketches. *International Conference on Machine Learning*, pages 166–175, 2017.
- Bowen Baker, Ingmar Kanitscheider, Todor Markov, Yi Wu, Glenn Powell, Bob McGrew, and Igor Mordatch. Emergent tool use from multi-agent autotutorials. *International Conference on Learning Representations*, 2019.

- Bowen Baker, Ilge Akkaya, Peter Zhokhov, Joost Huizinga, Jie Tang, Adrien Ecoffet, Brandon Houghton, Raul Sampedro, and Jeff Clune. Video pretraining (vpt): Learning to act by watching unlabeled online videos. *Advances in Neural Information Processing Systems*, 35:24639–24654, 2022.
- Anton Bakhtin, Noam Brown, Emily Dinan, Gabriele Farina, Colin Flaherty, Daniel Fried, Andrew Goff, Jonathan Gray, Hengyuan Hu, Athul Paul Jacob, et al. Human-level play in the game of diplomacy by combining language models with strategic reasoning. *Science*, 378(6624):1067–1074, 2022.
- Nolan Bard, Jakob N Foerster, Sarath Chandar, Neil Burch, Marc Lanctot, H Francis Song, Emilio Parisotto, Vincent Dumoulin, Subhodeep Moitra, Edward Hughes, et al. The hanabi challenge: A new frontier for ai research. *Artificial Intelligence*, 280:103216, 2020.
- Peter W Battaglia, Jessica B Hamrick, Victor Bapst, Alvaro Sanchez-Gonzalez, Vinicius Zambaldi, Mateusz Malinowski, Andrea Tacchetti, David Raposo, Adam Santoro, Ryan Faulkner, et al. Relational inductive biases, deep learning, and graph networks. *arXiv preprint arXiv:1806.01261*, 2018.
- Charles Beattie, Joel Z Leibo, Denis Teplyashin, Tom Ward, Marcus Wainwright, Heinrich Küttler, Andrew Lefrancq, Tim Green, Víctor Valdés, Amir Sadik, et al. Deepmind lab. *arXiv preprint arXiv:1612.03801*, 2016.
- Marc G Bellemare, Yavar Naddaf, Joel Veness, and Michael Bowling. The arcade learning environment: An evaluation platform for general agents. *Journal of Artificial Intelligence Research*, 47:253–279, 2013.
- Richard Bellman. A markovian decision process. *Journal of mathematics and mechanics*, pages 679–684, 1957.
- Christopher Berner, Greg Brockman, Brooke Chan, Vicki Cheung, Przemysław Debiak, Christy Dennison, David Farhi, Quirin Fischer, Shariq Hashme, Chris Hesse, et al. Dota 2 with large scale deep reinforcement learning. *arXiv preprint arXiv:1912.06680*, 2019.
- Michael Bowling and Manuela Veloso. Multiagent learning using a variable learning rate. *Artificial Intelligence*, 136(2): 215–250, 2002.
- Noam Brown and Tuomas Sandholm. Superhuman ai for multiplayer poker. *Science*, 365(6456):885–890, 2019.
- Lucian Busoniu, Robert Babuska, and Bart De Schutter. A comprehensive survey of multiagent reinforcement learning. *IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews)*, 38(2):156–172, 2008.
- Colin F Camerer. *Behavioral game theory: Experiments in strategic interaction*. Princeton University Press, 2004.
- Murray Campbell, A Joseph Hoane, and Feng-hsiung Hsu. Deep blue. *Artificial Intelligence*, 134(1-2):57–83, 2002.
- Micah Carroll, Rohin Shah, Mark K Ho, Thomas L Griffiths, Sanjit A Seshia, Pieter Abbeel, and Anca Dragan. On the utility of learning about humans for human-ai coordination. *Advances in Neural Information Processing Systems*, 32:5174–5185, 2019.
- Yu-Han Chang, Tracey Ho, and Leslie Pack Kaelbling. All learning is local: Multi-agent learning in global reward games. In *Advances in Neural Information Processing Systems*, pages 807–814, 2003.
- Yu Chen, Michael Thompson, and Carlos Garcia. Opponent modeling and deck adaptation in collectible card games. *International Joint Conference on Artificial Intelligence*, pages 234–241, 2023.
- Yu Chen, Xiaoming Wang, and Wei Zhang. Emergent social norms in multi-agent minecraft societies. *International Conference on Autonomous Agents and MultiAgent Systems*, pages 345–353, 2024.
- Yinlam Chow, Ofir Nachum, Edgar Duenez-Guzman, and Mohammad Ghavamzadeh. A lyapunov-based approach to safe reinforcement learning. *Advances in Neural Information Processing Systems*, 31:8103–8112, 2018.
- Karl Cobbe, Oleg Klimov, Chris Hesse, Taehoon Kim, and John Schulman. Quantifying generalization in reinforcement learning. *International Conference on Machine Learning*, pages 1282–1289, 2019.
- Cédric Colas, Olivier Sigaud, and Pierre-Yves Oudeyer. How many random seeds? statistical power analysis in deep reinforcement learning experiments. *arXiv preprint arXiv:1806.08295*, 2018.
- Wojciech M Czarnecki, Gauthier Gidel, Brendan Tracey, Karl Tuyls, Shayegan Omidshafiei, David Balduzzi, and Max Jaderberg. A framework for population-based training. *Advances in Neural Information Processing Systems*, 33:18157–18168, 2020.
- Abhishek Das, Theophile Gervet, Joshua Romoff, Dhruv Batra, Devi Parikh, Mike Rabbat, and Joelle Pineau. Tarmac: Targeted multi-agent communication. *International Conference on Machine Learning*, pages 1538–1546, 2019.

- Alexey Dosovitskiy and Vladlen Koltun. Learning to act by watching unlabeled videos. *arXiv preprint arXiv:1712.01956*, 2017.
- Lasse Espeholt, Hubert Soyer, Remi Munos, Karen Simonyan, Vlad Mnih, Tom Ward, Yotam Doron, Vlad Firoiu, Tim Harley, Iain Dunning, et al. Impala: Scalable distributed deep-rl with importance weighted actor-learner architectures. *International Conference on Machine Learning*, pages 1407–1416, 2018.
- Katrina Evtimova, Andrew Drozdov, Douwe Kiela, and Kyunghyun Cho. Emergent communication in a multi-modal, multi-step referential game. *International Conference on Learning Representations*, 2018.
- Lixin Fan, Guanzhi Wang, Yunfan Jiang, Ajay Mandlekar, Yuke Zhu, and Jie Tang. Language-guided multi-agent coordination in open-world environments. *International Conference on Machine Learning*, 2024.
- Jesse Farebrother, Marlos C Machado, and Michael Bowling. Generalization and regularization in dqn. *arXiv preprint arXiv:1810.00123*, 2018.
- Jakob Foerster, Nantas Nardelli, Gregory Farquhar, Triantafyllos Afouras, Philip HS Torr, Shimon Whiteson, et al. Stabilising experience replay for deep multi-agent reinforcement learning. *International Conference on Machine Learning*, pages 1146–1155, 2017.
- Jakob Foerster, Gregory Farquhar, Triantafyllos Afouras, Nantas Nardelli, and Shimon Whiteson. Counterfactual multi-agent policy gradients. *Proceedings of the AAAI Conference on Artificial Intelligence*, 32(1):2974–2982, 2018.
- Jakob Foerster, Francis Song, Edward Hughes, Neil Burch, Iain Dunning, Shimon Whiteson, Matthew Botvinick, and Michael Bowling. Bayesian action decoder for deep multi-agent reinforcement learning. *International Conference on Machine Learning*, pages 1942–1951, 2019.
- Javier García and Fernando Fernández. A comprehensive survey on safe reinforcement learning. *Journal of Machine Learning Research*, 16(1):1437–1480, 2015.
- Adam Gleave, Michael Dennis, Charles Wild, Neel Kant, Sergey Levine, and Stuart Russell. Adversarial attacks on multi-agent reinforcement learning. *International Conference on Machine Learning*, pages 3649–3659, 2020.
- Jonathan Gray, Adam Lerer, Anton Bakhtin, and Noam Brown. Learning to communicate and negotiate in imperfect information games. *Advances in Neural Information Processing Systems*, 33:1234–1245, 2020.
- Sam Greydanus, Anurag Koul, Jonathan Dodge, and Alan Fern. Visualizing and understanding atari agents. *International Conference on Machine Learning*, pages 1792–1801, 2018.
- Aditya Grover, Maruan Al-Shedivat, Jayesh Gupta, Yuri Burda, and Harrison Edwards. Best-response learning in multi-agent games with unknown rewards. *Advances in Neural Information Processing Systems*, 31:10256–10267, 2018.
- William H Guss, Brandon Houghton, Nicholay Topin, Phillip Wang, Cayden Codel, Manuela Veloso, and Ruslan Salakhutdinov. Minerl: A large-scale dataset of minecraft demonstrations. *Advances in Neural Information Processing Systems*, 32:2292–2300, 2019.
- Dylan Hadfield-Menell, Stuart J Russell, Pieter Abbeel, and Anca Dragan. Cooperative inverse reinforcement learning. *Advances in Neural Information Processing Systems*, 29:3916–3924, 2016.
- Eric A Hansen, Daniel S Bernstein, and Shlomo Zilberstein. Dynamic programming for partially observable stochastic games. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 709–715, 2004.
- Matthew Hausknecht and Peter Stone. Deep recurrent q-learning for partially observable mdps. *arXiv preprint arXiv:1507.06527*, 2015.
- Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. Deep reinforcement learning that matters. *Proceedings of the AAAI Conference on Artificial Intelligence*, 32(1):3207–3214, 2018.
- Pablo Hernandez-Leal, Bilal Kartal, and Matthew E Taylor. A survey and critique of multiagent deep reinforcement learning. *Autonomous Agents and Multi-Agent Systems*, 33(6):750–797, 2019.
- Edward Hughes, Joel Z Leibo, Matthew Phillips, Karl Tuyls, Edgar Dueñez-Guzman, Antonio García Castañeda, Iain Dunning, Tina Zhu, Kevin McKee, Raphael Koster, et al. Inequity aversion improves cooperation in intertemporal social dilemmas. *Advances in Neural Information Processing Systems*, 31:3330–3340, 2018.
- Maximilian Igl, Kamil Ciosek, Yingzhen Li, Sebastian Tschiatschek, Cheng Zhang, Sam Devlin, and Katja Hofmann. Generalization in reinforcement learning with selective noise injection and information bottleneck. *Advances in Neural Information Processing Systems*, 32:13978–13990, 2019.

- Max Jaderberg, Valentin Dalibard, Simon Osindero, Wojciech M Czarnecki, Jeff Donahue, Ali Razavi, Oriol Vinyals, Tim Green, Iain Dunning, Karen Simonyan, et al. Population based training of neural networks. *arXiv preprint arXiv:1711.09846*, 2017.
- Max Jaderberg, Wojciech M Czarnecki, Iain Dunning, Luke Marris, Guy Lever, Antonio Garcia Castaneda, Charles Beattie, Neil C Rabinowitz, Ari S Morcos, Avraham Ruderman, et al. Human-level performance in first-person multiplayer games with population-based deep reinforcement learning. *Science*, 364(6443):859–865, 2019.
- Natasha Jaques, Angeliki Lazaridou, Edward Hughes, Caglar Gulcehre, Pedro A Ortega, DJ Strouse, Joel Z Leibo, and Nando de Freitas. Social influence as intrinsic motivation for multi-agent deep reinforcement learning. *International Conference on Machine Learning*, pages 3040–3049, 2019.
- Matthew Johnson, Katja Hofmann, Tim Hutton, and David Bignell. The malmo platform for artificial intelligence experimentation. *International Joint Conference on Artificial Intelligence*, pages 4246–4247, 2016.
- Niels Justesen, Ruben Rodriguez Torrado, Philip Bontrager, Ahmed Khalifa, Julian Togelius, and Sebastian Risi. Learning to generalize in reinforcement learning. *IEEE Transactions on Games*, 11(3):219–230, 2018.
- Michał Kempka, Marek Wydmuch, Grzegorz Runc, Jakub Toczek, and Wojciech Jaśkowski. Vizdoom: A doom-based ai research platform for visual reinforcement learning. *arXiv preprint arXiv:1605.02097*, 2016.
- Dong-Ki Kim, Miao Liu, Matthew Riemer, Mingfei Sun, Hany Abdulsamad, and Gerald Tesauro. Learning to coordinate with meta-learned policies. *International Conference on Autonomous Agents and MultiAgent Systems*, pages 712–720, 2021.
- Woojun Kim, Jongeui Park, and Youngchul Sung. Learning hierarchical communication in multi-agent reinforcement learning. *International Conference on Learning Representations*, 2020.
- Jakub Kuba, Rui Chen, Muning Wen, Ying Wen, Jun Zhang, Jun Wang, and Yaodong Yang. A unified framework for multi-agent policy gradients. *Advances in Neural Information Processing Systems*, 35:31562–31575, 2022.
- Guillaume Lample and Devendra Singh Chaplot. Playing fps games with deep reinforcement learning. *Proceedings of the AAAI Conference on Artificial Intelligence*, 31(1):2140–2146, 2017.
- Marc Lanctot, Vinicius Zambaldi, Audrunas Gruslys, Angeliki Lazaridou, Karl Tuyls, Julien Pérolat, David Silver, and Thore Graepel. A unified game-theoretic approach to multiagent reinforcement learning. *Advances in Neural Information Processing Systems*, 30:4190–4203, 2017.
- Angeliki Lazaridou, Alexander Peysakhovich, and Marco Baroni. Multi-agent cooperation and the emergence of (natural) language. *International Conference on Learning Representations*, 2017.
- Joel Z Leibo, Vinicius Zambaldi, Marc Lanctot, Janusz Marecki, and Thore Graepel. Multi-agent reinforcement learning in sequential social dilemmas. *International Conference on Autonomous Agents and MultiAgent Systems*, pages 464–473, 2017.
- Stefanos Leonardos and Georgios Piliouras. The dynamics of gradient-based learning in two-player zero-sum games. *IEEE Transactions on Automatic Control*, 66(11):5117–5132, 2021.
- Michael L Littman. Markov games as a framework for multi-agent reinforcement learning. *Machine learning proceedings 1994*, pages 157–163, 1994.
- Ryan Lowe, Yi Wu, Aviv Tamar, Jean Harb, Pieter Abbeel, and Igor Mordatch. Multi-agent actor-critic for mixed cooperative-competitive environments. *Advances in Neural Information Processing Systems*, 30:6379–6390, 2017.
- Ryan Lowe, Jakob Foerster, Y-Lan Boureau, Joelle Pineau, and Yann Dauphin. On the emergence of compositional languages in multi-agent systems. *International Conference on Machine Learning*, pages 6355–6365, 2020.
- Xueguang Lyu, Yuchen Xiao, Brett Daley, and Christopher Amato. A unified view of centralized training with decentralized execution in multi-agent reinforcement learning. *arXiv preprint arXiv:2105.08578*, 2021.
- Anuj Mahajan, Tabish Rashid, Mikayel Samvelyan, and Shimon Whiteson. On the representational capacity of neural networks for multi-agent q-learning. *arXiv preprint arXiv:2102.12447*, 2021.
- Luke Marris, Paul Muller, Marc Lanctot, Karl Tuyls, and Thore Graepel. Multi-agent training beyond zero-sum. *International Conference on Machine Learning*, pages 15097–15113, 2022.
- Stephen McAleer, John Lanier, Roy Fox, and Pierre Baldi. Pipeline psro: A scalable approach for finding approximate nash equilibria in large games. *Advances in Neural Information Processing Systems*, 34:20271–20283, 2021.

- Kevin McKee, Ian Gemp, Brian McWilliams, Edgar Duéñez-Guzmán, Edward Hughes, and Joel Z Leibo. Social diversity and social preferences in mixed-motive reinforcement learning. *International Conference on Autonomous Agents and MultiAgent Systems*, pages 869–877, 2020.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller. Playing atari with deep reinforcement learning. *arXiv preprint arXiv:1312.5602*, 2013.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 2015.
- Igor Mordatch and Pieter Abbeel. Emergence of grounded compositional language in multi-agent populations. *Proceedings of the AAAI Conference on Artificial Intelligence*, 32(1):5495–5502, 2018.
- Paul Muller, Shayegan Omidshafiei, Mark Rowland, Karl Tuyls, Julien Perolat, Siqi Liu, Daniel Hennes, Luke Marris, Marc Lanctot, Edward Hughes, et al. A generalized training approach for multiagent learning. *International Conference on Learning Representations*, 2020.
- Andrew Y Ng and Stuart J Russell. Algorithms for inverse reinforcement learning. *International Conference on Machine Learning*, pages 663–670, 2000.
- Santiago Ontañón, Gabriel Synnaeve, Alberto Uriarte, Florian Richoux, David Churchill, and Mike Preuss. A survey of real-time strategy game ai research and competition in starcraft. *IEEE Transactions on Computational Intelligence and AI in Games*, 5(4):293–311, 2013.
- Gregory Palmer, Karl Tuyls, Daan Bloembergen, and Rahul Savani. Lenient multi-agent deep reinforcement learning. *International Conference on Autonomous Agents and MultiAgent Systems*, pages 443–451, 2018.
- Georgios Papoudakis, Filippos Christoffersen, and Stefano V Albrecht. Benchmarking multi-agent deep reinforcement learning algorithms in cooperative tasks. *arXiv preprint arXiv:2006.07869*, 2020.
- Joon Sung Park, Joseph O’Brien, Carrie Cai, Meredith Morris, Percy Liang, and Michael Bernstein. Generative agents: Interactive simulacra of human behavior. *Proceedings of the ACM Symposium on User Interface Software and Technology*, pages 1–22, 2023.
- Tim Pearce, Alex Zhu, Kianté Brantley, and Nan Jiang. Counter-strike: A benchmark for cooperative-competitive multi-agent learning. *Advances in Neural Information Processing Systems*, 36:4567–4580, 2023.
- Julien Pérolat, Bruno Scherrer, Bilal Piot, and Olivier Pietquin. On the convergence of reinforcement learning in general-sum stochastic games. *arXiv preprint arXiv:1507.00387*, 2015.
- Alexander Peysakhovich and Adam Lerer. Prosocial learning agents solve generalized stag hunts better than selfish ones. *Proceedings of the International Joint Conference on Autonomous Agents and Multiagent Systems*, pages 2043–2044, 2018.
- Jonathan Rabiner, Nolan Bard, and Jakob Foerster. Understanding cooperative behavior in hanabi through behavioral analysis. *arXiv preprint arXiv:2103.08215*, 2021.
- Neil C Rabinowitz, Frank Perbet, H Francis Song, Chiyuan Zhang, SM Ali Eslami, and Matthew Botvinick. Machine theory of mind. *International Conference on Machine Learning*, pages 4218–4227, 2018.
- Roberta Raileanu, Emily Denton, Arthur Szlam, and Rob Fergus. Modeling others using oneself in multi-agent reinforcement learning. *International Conference on Machine Learning*, pages 4257–4266, 2018.
- Tabish Rashid, Mikayel Samvelyan, Christian Schroeder De Witt, Gregory Farquhar, Jakob Foerster, and Shimon Whiteson. Qmix: Monotonic value function factorisation for deep multi-agent reinforcement learning. In *International Conference on Machine Learning*, pages 4295–4304, 2018.
- Tabish Rashid, Gregory Farquhar, Bei Peng, and Shimon Whiteson. Weighted qmix: Expanding monotonic value function factorisation. *arXiv preprint arXiv:2006.10800*, 2020.
- Mikayel Samvelyan, Tabish Rashid, Christian Schroeder De Witt, Gregory Farquhar, Nantas Nardelli, Tim GJ Rudner, Chia-Man Hung, Philip HS Torr, Jakob Foerster, and Shimon Whiteson. The starcraft multi-agent challenge. *arXiv preprint arXiv:1902.04043*, 2019.
- Christian Schroeder de Witt, Tarun Gupta, Denys Makoviichuk, Viktor Makovychuk, Philip HS Torr, Mingfei Sun, and Shimon Whiteson. Eligibility traces for multi-agent reinforcement learning. *International Conference on Machine Learning*, pages 8532–8542, 2020.

- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Lloyd S Shapley. Stochastic games. *Proceedings of the national academy of sciences*, 39(10):1095–1100, 1953.
- David Silver, Aja Huang, Chris J Maddison, Arthur Guez, Laurent Sifre, George Van Den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, et al. Mastering the game of go with deep neural networks and tree search. *Nature*, 529(7587):484–489, 2016.
- David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, et al. Mastering the game of go without human knowledge. *Nature*, 550(7676):354–359, 2017.
- Amanpreet Singh, Tushar Jain, and Sainbayar Sukhbaatar. Learning when to communicate at scale in multiagent cooperative and competitive tasks. *International Conference on Learning Representations*, 2019.
- Jennifer Smith, Liang Chen, and Rajesh Kumar. Emergent communication through in-game actions in first-person shooter environments. *International Conference on Learning Representations*, 2024.
- Kyunghwan Son, Daewoo Kim, Wan Ju Kang, David Hostallero, and Yung Yi. Qtran: Learning to factorize with transformation for cooperative multi-agent reinforcement learning. In *International Conference on Machine Learning*, pages 5887–5896, 2019.
- Joseph Suarez, Yilun Du, Phillip Zhu, and Igor Mordatch. Neural mmo: A massively multiagent game environment for training and evaluating intelligent agents. *arXiv preprint arXiv:1903.00784*, 2019.
- Sainbayar Sukhbaatar, Arthur Szlam, and Rob Fergus. Learning multiagent communication with backpropagation. *Advances in Neural Information Processing Systems*, 29:2244–2252, 2016.
- Peter Sunehag, Guy Lever, Audrunas Gruslys, Wojciech Marian Czarnecki, Vinicius Zambaldi, Max Jaderberg, Marc Lanctot, Nicolas Sonnerat, Joel Z Leibo, Karl Tuyls, et al. Value-decomposition networks for cooperative multi-agent reinforcement learning. *arXiv preprint arXiv:1706.05296*, 2017.
- Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*. MIT press, 1998.
- Richard S Sutton, David McAllester, Satinder Singh, and Yishay Mansour. Policy gradient methods for reinforcement learning with function approximation. *Advances in Neural Information Processing Systems*, 12:1057–1063, 2000.
- Hongyao Tang, Jianye Hao, Tangjie Lv, Yingfeng Chen, Zongzhang Zhang, Hangtian Jia, Chunxu Ren, Yan Zheng, Changjie Fan, and Li Wang. Hierarchical deep multiagent reinforcement learning with temporal abstraction. *arXiv preprint arXiv:1809.09332*, 2018.
- Josh Tobin, Rachel Fong, Alex Ray, Jonas Schneider, Wojciech Zaremba, and Pieter Abbeel. Domain randomization for transferring deep neural networks from simulation to the real world. *IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 23–30, 2017.
- Alexander S Vezhnevets, Simon Osindero, Tom Schaul, Nicolas Heess, Max Jaderberg, David Silver, and Koray Kavukcuoglu. Feudal networks for hierarchical reinforcement learning. *International Conference on Machine Learning*, pages 3540–3549, 2017.
- Oriol Vinyals, Igor Babuschkin, Wojciech M Czarnecki, Michaël Mathieu, Andrew Dudzik, Junyoung Chung, David H Choi, Richard Powell, Timo Ewalds, Petko Georgiev, et al. Grandmaster level in starcraft ii using multi-agent reinforcement learning. *Nature*, 575(7782):350–354, 2019.
- Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Lixin Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023.
- Jianhao Wang, Zhizhou Ren, Terry Liu, Yang Yu, and Chongjie Zhang. Qplex: Duplex dueling multi-agent q-learning. *International Conference on Learning Representations*, 2021.
- Tonghan Wang, Tarun Gupta, Anuj Mahajan, Bei Peng, Shimon Whiteson, and Chongjie Zhang. Counterfactual multi-agent policy gradients with value decomposition. *Advances in Neural Information Processing Systems*, 33:5678–5689, 2020.
- Christopher JCH Watkins and Peter Dayan. Q-learning. *Machine learning*, 8(3-4):279–292, 1992.
- Muning Wen, Jakub Kuba, Runji Lin, Weinan Zhang, Ying Wen, Jun Wang, and Yaodong Yang. Multi-agent reinforcement learning with transformer networks. *International Conference on Learning Representations*, 2022.
- David H Wolpert and Kagan Tumer. Optimal payoff functions for members of collectives. *Advances in Complex Systems*, 4(02n03):265–279, 2001.

- Hao Wu, Deheng Ye, Zhao Liu, Mingfei Sun, Bei Shi, Shikun Yang, and Guojun Li. Multi-agent reinforcement learning for team-based moba games with attention mechanisms. *IEEE Transactions on Games*, 14(2):234–245, 2021.
- Shikun Yang, Deheng Ye, Zhao Liu, Hao Wu, Mingfei Sun, Bei Shi, and Guojun Li. Counterfactual credit assignment in moba games. *IEEE Transactions on Games*, 15(3):456–468, 2022.
- Yaodong Yang, Rui Luo, Minne Li, Ming Zhou, Weinan Zhang, and Jun Wang. Mean field multi-agent reinforcement learning. *International Conference on Machine Learning*, pages 5571–5580, 2018.
- Deheng Ye, Zhao Liu, Mingfei Sun, Bei Shi, Peilin Zhao, Hao Wu, Hongsheng Yu, Shikun Yang, Xipeng Wu, Qingwei Guo, et al. Towards playing full moba games with deep reinforcement learning. *Advances in Neural Information Processing Systems*, 33:621–632, 2020.
- Chao Yu, Akash Velu, Eugene Vinytsky, Jiaxuan Gao, Yu Wang, Alexandre Bayen, and Yi Wu. The surprising effectiveness of ppo in cooperative multi-agent games. *Advances in Neural Information Processing Systems*, 35:24611–24624, 2022.
- Shusen Zhang, Tong Liu, Jun Wang, and Wenyin Liu. A survey on deep learning for steering in games. *IEEE Transactions on Games*, 11(3):205–218, 2018.
- Tianjun Zhang, Huazhe Xu, Xiaolong Wang, Yi Wu, Kurt Keutzer, Joseph E Gonzalez, and Yuandong Tian. Multi-agent world models: Learning and planning in multi-agent environments. *Advances in Neural Information Processing Systems*, 34:1234–1246, 2021.
- Yichi Zhang, Yang Yu, and Chongjie Zhang. Learning and reasoning in multi-agent systems: A level-k perspective. *International Conference on Autonomous Agents and MultiAgent Systems*, pages 1624–1632, 2020.
- Yiming Zhang, Wei Liu, and Xiaojing Wang. Multi-agent reinforcement learning for auto-battler games. *IEEE Conference on Games*, pages 123–130, 2022.
- Yiming Zhang, Yu Chen, and Michael Thompson. Large language models as negotiators in multi-agent games. *International Conference on Learning Representations*, 2024.